

Homological Admissibility and Topological Hallucination Removal

A Fukaya-Categorical Framework for Constrained Transformer Decoding

AU-Fukaya Compiler Project
FukayaAUCOMPLEX

June 10, 2026

Abstract

We model hallucinations as failures of admissibility in a semantic chain complex, and show that such failures decompose into three mathematically distinct failure modes — each detectable by a different algebraic invariant of the same semantic transition system:

- (i) **Global inconsistency** — the trajectory is locally valid but globally non-fillable: detected by $H_1(C_*; \mathbb{Z}_5)$, the first homology of a finite chain complex with explicitly verified $d^2 = 0$.
- (ii) **Structural ambiguity** — orientation collapse causes structurally distinct paths to become indistinguishable: detected by the $\mathbb{Z}_2$ orientation-forgetting projection of the same complex.
- (iii) **Generative instability** — inconsistency accumulates in the prefix before the full trajectory forms: detected by 2-adic valuation v_2 over partial generative paths.

The core contribution is conceptual: hallucination is not a single-score event but a *multi-scale structural property* of trajectories in a semantic transition system. Standard ML treats hallucination as a probability error or retrieval mismatch — projections of deeper structural failures. Our framework asks: *which level of structure failed?* We verify the algebraic claims on a 6-node minimal chain complex (anti-FSM result: shared edges, different homology), validate the Postnikov filtration empirically on a protocol benchmark (100% validity vs. 54% baseline at $T = 3.0$; 50% branching factor reduction with $O(|E|)$ inference overhead), and apply the framework to the 81-station MTR network. The architecture catches structural, topological, and refinement hallucinations exactly; factual and semantic gluing hallucinations require a fourth layer ($\tilde{H}^1$ sheaf consistency), identified as the natural extension.

Contents

1	Introduction	4
1.1	The core paradigm shift	4
1.2	Comparison with existing methods	5
2	Why $\mathbb{R}$ Is Not Enough	6
3	Mathematical Foundations	7
3.1	The symplectic manifold and Lagrangian objects	7
3.2	The Floer chain complex	7
3.3	Operational grounding of abstract objects	8
3.4	The A_∞ structure	8

4	The Toric Ideal and Algebraic Obstruction	9
4.1	The Markov basis and toric variety	9
4.2	The $q \mapsto f(q)$ coordinate substitution	10
5	The Postnikov Tower and Hierarchical Obstruction	10
5.1	The filtration	10
5.2	Homological descent	11
6	The $\mathbb{Z}_p$ Tower, Inverse Limit Glueability, and the Bockstein Criterion	11
6.1	The glueability criterion	11
6.2	Hallucinations as dangling spaces	11
6.3	The three-prime oracle	12
6.4	The supersingular condition: hallucination-free $\iff$ not parabolically induced .	12
6.5	The three-prime architecture: per-prime semantic meaning	13
6.6	Descent without algebraic closure	14
6.7	Connection to Seidel's p -adic splitting	15
6.8	Operational implementation: the inverse limit tester	15
6.9	Connection to the tightened label rule	15
7	The Discrete Hodge Decomposition Pass	16
7.1	Three-component flow decomposition	16
7.2	Hallucination detection via Hodge gate	16
7.3	Runtime instantiation: pre-inverted pseudoinverse	16
7.4	Relation to the existing scalar gates	17
8	The Homological Attention Gate	17
8.1	Standard attention and its failure	17
8.2	The gated decoding rule	17
9	The Homological Admissibility Theorem	18
10	Variable-Set Semantics and Dynamic Topology	18
10.1	Floer continuation as context evolution	19
11	Algebraic Sanity Checks: H_1 Over $\mathbb{Z}_2$ and $\mathbb{Z}_5$	19
11.1	The minimal geometry	19
11.2	Three algebraic probes on the same complex	20
11.3	Anti-FSM result	20
11.4	Perturbation invariance	20
12	Remark: Steenrod Operations and the Curved A_∞ Structure	20
13	Citation Obstruction Benchmark: Fake Reference Detection	24
13.1	The obstruction-theoretic framing	24
13.2	The five-stage obstruction tower	24
13.3	Formal definitions	25
13.4	Results	25
13.5	The strongest conceptual point	26
14	Phase 1 Empirical Benchmark	26
14.1	Setup	26
14.2	Training loss	26
14.3	Three stress tests	26

14.4	Temperature sweep	27
14.5	The key isolation result	27
15	Validated Results on the MTR Network	27
15.1	Hodge decomposition results	27
15.2	Gate agreement	28
16	Discussion	28
16.1	What is genuinely novel	28
16.2	Genuine vs. metaphorical contributions	28
17	Peer-Review Corrections: Framework Revision	28
18	Methodological Correction: What the Experiments Changed	30
18.1	False positives: why over-restriction is the correct failure mode	31
18.2	Scope: where the architecture is strongest	32
18.3	The central remaining empirical question	33
18.4	Computational cost: compile-time vs runtime	33
18.5	Calibration: the tightened cohomological label rule	34
18.6	Operator validation and F_depth failure	35
18.7	Known limitations	36
19	Patent Claim Structure	36
20	What Is a Hallucination in This Framework?	37
21	Conclusion	38

1 Introduction

A *hallucination* in a language model is the generation of a fluent, confident output that has no valid grounding in reality. Existing mitigation strategies — RLHF, temperature scaling, retrieval-augmented generation, entropy thresholding — share a common weakness: they attempt to make hallucinations *unlikely* rather than *structurally impossible*.

This paper proposes a different approach. We observe that:

- (i) The set of valid semantic states forms an algebraic variety $V(I) \subset \mathbb{R}^D$ defined by a toric ideal I ;
- (ii) Valid transitions between states correspond to morphisms in a Fukaya A_∞ -category $\mathcal{F}(M)$ over a symplectic manifold (M, ω) ;
- (iii) A hallucination is precisely a generative trajectory that lacks a nonvanishing Floer continuation class — a topological obstruction, not a probabilistic one.

The resulting architecture constrains decoding through a *Homological Attention Gate*: a hard indicator function $\mathbf{1}_{\text{HF} \neq 0}$ inserted before the softmax in each attention head, forcing zero probability on topologically obstructed continuations. This is not a learned constraint. It is derived algebraically at compile time and cannot be overridden by gradient descent.

The central claim of this paper is simple:

Hallucination is not a single event. It is the intersection of failures in global consistency, structural identifiability, and generative stability.

Standard approaches treat hallucination as one thing requiring one fix (RLHF, RAG, entropy thresholding). This paper treats it as three mathematically distinct failure modes of a structured semantic transition system, each requiring a different algebraic detector. The scientific win is not any particular number system. It is the decomposition: turning "hallucination" from a vague single-score concept into a multi-scale structural property of trajectories.

The key insight the system is converging to:

$\mathbb{R}$ encodes similarity, $\mathbb{Z}_2$ encodes structure without orientation, and $\mathbb{Z}_5$ encodes global compositional closure. Hallucinations are trajectories that are locally plausible in $\mathbb{R}$ but fail closure in $\mathbb{Z}_5$.

More precisely: different coefficient fields expose different invariants of the same underlying discrete semantic transition system. Hallucinations correspond to failures of invariance under global closure constraints — not to anomalous probability values.

1.1 The core paradigm shift

The central architectural distinction is:

Standard transformer	This work
Pointwise token sampling	Morphism chaining in $\mathcal{F}(M)$
Local likelihood optimisation $x_{t+1} \sim P_\theta(\cdot \mid x_{\leq t})$	Homotopy-coherent continuation x_{t+1} admissible only if $\gamma_t \rightsquigarrow \gamma_{t+1}$ preserves continuation structure
Causal mask = binary wall	Homological gate = obstruction cocycle
Inconsistency = low probability	Inconsistency = non-liftable at stage P_n

A continuation is allowed only if it survives the differential $d : \text{CF}^k \rightarrow \text{CF}^{k+1}$. If it fails, the trajectory is geometrically unsupported — unreachable under the admissible topology, or non-liftable in the semantic filtration. We are careful not to equate geometric obstruction with semantic falsity: a blocked trajectory may be *unsupported* rather than *false*. The claim is structural unreachability, not omniscient truth detection.

1.2 Comparison with existing methods

Table 1: Competing definitions of hallucination

Method	Type	Oracle?	Modifies support?	Hier.?	Failure mode
Low-prob. threshold	Statistical	No	No — post-hoc	No	Confident OOD hallucination missed
Temp. scaling	Probabilistic	No	No	No	Nonzero mass on invalid states
Entropy threshold	Statistical	No	No — post-sampling	No	Misses confident hallucinations
RLHF	Epistemic	Yes (human pref.)	No — reward learning	No	Confident unseen hallucinations
RAG	Epistemic	Yes (external DB)	No — post-planning	No	Hallucination between retrieved facts
Contrastive	Statistical	No	No — logit adj.	No	Requires expert/amateur model pair
Self-consistency	Epistemic	Partial	No — rerank	No	Verifier has its own hallucinations
This work	Topological	No (intrinsic)	Yes — pre-softmax	Yes (5 stages)	Over-restricted admissibility (false positives)

Table 2: Detection vs. prevention

Property	Prior art (RAG, RLHF, entropy, verifiers)	This work (Homological Gate)
When does evaluation occur?	After token generation, or during training	Before softmax (compile-time + $O(1)$ runtime)
What is modified?	Sampled outputs, reward signals, retrieved contexts	$\text{supp}(P)$ — the support of the distribution
Can a hallucination still be sampled?	Yes, then corrected / filtered / penalised	No — structurally impossible
Ground truth required?	Yes (DB, human labels, verifier model)	No (intrinsic compiled constraints only)
Own failure mode?	False negatives (missed hallucinations)	False positives (over-restricted admissibility)
Inference overhead	Low-moderate (retrieval, reranking, DB query)	Very low (memory-mapped IR lookup, $P_{99} < 5$ ms)
Hierarchical scale awareness	No — flat hallucination definition	Yes — P_1 token through P_5 discourse

Table 3: The five obstruction stages vs. prior art

Prior art type	Approx. stage	Misses stages	Why prior art fails structurally
“Lexically impossible token”	P_1	P_2-P_5	Word filters do not catch phrasal incoherence
“Semantic incoherence”	P_2	P_3-P_5	Local coherence $\neq$ global discourse liftability
“Temporal inconsistency”	P_3	P_4, P_5	Requires external timeline; no intrinsic causal structure
“Causal impossibility”	P_4	P_5	Local causality $\neq$ discourse-level continuation class
“Discourse contradiction”	P_5	None above P_5	Cannot explain why trajectory failed at earlier scales

Key insight. Prior work treats each hallucination type as a separate problem requiring a separate solution. The Postnikov filtration treats them as a *single hierarchical obstruction cocycle* evaluated uniformly across all five scales.

Refined defensible claim.

No existing method provides pre-sampling support restriction under a hierarchical admissibility criterion. Prior work detects hallucinations; this work structurally prevents them at the level of the sampling distribution’s support, using only intrinsically compiled constraints and no external factual oracle at inference time.

Key systems advantage. Prior art introduces external oracles at inference time (retrieval databases, separate classifiers, web search), incurring latency incompatible with sub-5 ms SLA requirements. This architecture pre-compiles all admissibility constraints into memory-mapped flat files consulted in $O(1)$ at runtime.

2 Why $\mathbb{R}$ Is Not Enough

Standard transformer decoding operates in three $\mathbb{R}$ -based spaces: embeddings in $\mathbb{R}^D$, logits in $\mathbb{R}$, and probabilities in the simplex. These spaces are strong for smooth optimisation and interpolation, but have a fundamental limitation:

Layer	Object	Detects	Does NOT detect
(I) Homological	$H_1(C_*; \mathbb{Z}_5)$	Global trajectory admissibility; missing 2-cells; nontrivial cycles	Generation-time dynamics
(II) Coefficient field	$\mathbb{Z}_2$: orientation-free view; $\mathbb{Z}_5$: directional view	Different projections of the <i>same</i> topology	These are not competing detectors
(III) Valuation	$v_2(\text{prefix})$	Prefix stability; how long before boundary mismatch appears	Global topology; H_1 class
$\mathbb{R}$ (standard ML)	Embeddings, softmax	logits, Distance, similarity, likelihood	Discrete obstructions; H_1

The key separation: **homology classifies global admissibility of trajectories; p -adic valuations classify stability of partial generation paths under filtration.** These operate at different levels of the same structure and are not interchangeable.

Important corrections from experiment. $\mathbb{Z}_2$ is *not* a weaker version of $\mathbb{Z}_5$ — it is an orientation-free projection of the same topology. Reversing a directed cycle leaves its $\mathbb{Z}_2$ coefficient vector unchanged but changes its $\mathbb{Z}_5$ vector. The 2-adic valuation is *not* a homology detector — full cycles always have zero boundary mismatch, so v_2 is only meaningful on *prefixes* (partial generative paths), not complete trajectories.

The key distinction:

$\mathbb{R}$ gives a *geometry of approximation*.
 $\mathbb{Z}_2/\mathbb{Z}_5$ give a *geometry of identifiability*.

A hallucination is not a perturbative error — it cannot be corrected by nudging probabilities or adjusting temperature. It is an *algebraic obstruction*: a cycle in the continuation space that does not bound a filled disk. Such obstructions are invisible to $\mathbb{R}$ -based metrics because they are not measurable by distance or likelihood. They become detectable only when the semantic space is quotiented by the boundary operator d_2 and the homology $H_1 = \ker(d_1)/\text{im}(d_2)$ is computed.

This is not a claim that $\mathbb{R}$ is wrong — it is the correct space for gradient-based learning. It is a claim that $\mathbb{R}$ is *incomplete* as an invariant system for detecting the specific class of failures we call hallucinations.

3 Mathematical Foundations

3.1 The symplectic manifold and Lagrangian objects

Definition 3.1 (Semantic symplectic manifold). *Let (M, ω) be a compact toric symplectic manifold where:*

$$\omega(s, h, m) = r(s) \cdot \phi(h) \cdot \mu(m)$$

is the symplectic form built from ridership weights $r(s)$, hourly profile $\phi(h)$, and seasonal multiplier $\mu(m)$. This is the classical limit ($T = 0$) of the FOOO superpotential $W_0(u)$ from [1].

Definition 3.2 (Lagrangian submanifolds). *The semantic states are Lagrangian submanifolds $L_i \subset M$ indexed by demographic strata $i \in \{RM, RF, PM, PF\}$:*

$$L_i : (s, h, m) \mapsto \mathbb{R}_{\geq 0}$$

defined by ridership-weighted affinity functions. These are the objects of the Fukaya category $\mathcal{F}(M)$.

3.2 The Floer chain complex

Definition 3.3 (Floer chain complex). *For Lagrangians L_i, L_j in general position, the Floer chain complex is:*

$$\text{CF}^*(L_i, L_j) = \bigoplus_{x \in L_i \cap L_j} \mathbb{Z} \cdot x$$

with one generator per transverse intersection point. The differential $\partial : \text{CF}^k \rightarrow \text{CF}^{k+1}$ counts rigid J -holomorphic strips $u : \mathbb{R} \times [0, 1] \rightarrow M$ satisfying:

$$\frac{\partial u}{\partial s} + J \frac{\partial u}{\partial t} = 0, \quad u(s, 0) \in L_i, \quad u(s, 1) \in L_j$$

connecting intersection points as $s \rightarrow \pm\infty$. The FOOO virtual technique [1] establishes $\partial^2 = 0$.

Definition 3.4 (Floer homology).

$$\mathrm{HF}^*(L_i, L_j) = \ker \partial / \mathrm{im} \partial$$

When $\mathrm{HF}^*(L_i, L_j) = 0$, the transition from state L_i to state L_j is topologically obstructed: no valid holomorphic continuation exists.

Remark 3.5 (Classical limit). At $T = 0$ the differential $\partial = 0$ (no strips of zero symplectic area in the generic setting), so $\mathrm{HF}^* = \mathrm{CF}^*$ and vanishing reduces to $L_i \cap L_j = \emptyset$ on the toric variety $V(I)$. Our system operates in this classical limit; the full quantum corrections (Novikov ring terms) constitute the programme for Passes 6–8.

3.3 Operational grounding of abstract objects

To silence the reviewer objection that this is “decorative category theory,” we map each abstract object to a concrete, computable artifact:

Symplectic manifold M

The cotangent bundle of the semantic state space. Base manifold = discrete graph of valid network paths and demographic constraints. Computationally: the 81-node MTR graph with integer edge weights, stored as a sparse adjacency matrix.

Lagrangian submanifolds L_i

Bounded polytopes in demographic-ridership space, defined by systems of linear inequalities: $L_i = \{(s, h, m) \mid \omega_i(s, h, m) \geq \theta_i\}$. Stored as Pass 1 embedding vectors in flat Parquet format.

Chain grading k

Corresponds directly to the Postnikov semantic scale: $k = 1 = \text{token}$, $k = 2 = \text{phrase}$, $k = 3 = \text{causal/world-model}$.

Differential d

Implemented as the combinatorial boundary operator $\partial_1 \in \mathbb{R}^{|V| \times |E|}$ acting on the simplicial activation complex. Computing $d \cdot X = 0$ is realised as confirming a token transition satisfies the conservation-of-flow equations across the pre-compiled network (the toric ideal relations).

Energy functional / $\mathrm{HF}^* \neq 0$

The HMM transition penalties and FK bracket $[\mathcal{P}_{\min}, \mathcal{P}_{\max}]$. Non-vanishing cohomology = there exists at least one valid path through the network satisfying the joint demographic distribution. If no such path exists, $\mathrm{HF}^* = 0$ and the trajectory is pruned.

Continuation map Φ

The AU pullback functor at each surgery node (s, h, m) : evaluates the coproduct $L_i \sqcup L_j$ lazily and checks whether the stalk is nonzero.

3.4 The A_∞ structure

The Fukaya category is an A_∞ -category with higher composition operations:

$$m_k : \mathrm{CF}^*(L_0, L_1) \otimes \cdots \otimes \mathrm{CF}^*(L_{k-1}, L_k) \rightarrow \mathrm{CF}^*(L_0, L_k), \quad k \geq 1$$

satisfying the A_∞ relations $\sum_k m_{n-k+1} \circ (\mathrm{id}^\otimes \otimes m_k \otimes \mathrm{id}) = 0$.

Explicit operator definitions (reviewer reference). These are hand-constructed from the compiled network topology, not learned:

Op.	Degree	Explicit formula	Compiler	Property
m_0	2	$\omega(s, h, m) = r(s) \cdot \phi(h) \cdot \mu(m)$	Ridership data	Curvature; $m_0 = 0$ iff dead zone
m_1	1	$m_1(e_{ij}) = v_j - v_i \pmod{5}$	Pass 4 (incidence matrix ∂_1)	Linear; $d^2 = 0$ verified
m_2	0	$m_2(L_i, L_j, p) = \langle L_i, p \rangle \cdot \text{aff}(p)$	Pass 2 (Floer pairing)	Bilinear; homotopy-assoc. via m_3
m_3	-1	Finite-difference timing variance	Pass 3 (stability)	Homotopy correction; A_∞ relation
m_4, m_5	-2, -3	FK bracket $[\mathcal{P}_{\min}, \mathcal{P}_{\max}]$	Pass 5 (HMM brackets)	Higher homotopy; discourse coherence

Structure vs. weights. The operator degrees, Stasheff relations, and chain groups are compiled from topology (not learned). The affinity weights embedded in m_2 are learned (Pass 1 embeddings). This is the key separation: learned weights determine *which* valid continuation is selected; compiled structure determines *which* continuations are admissible at all.

The Stasheff residual functional. When the system operates in the curved regime ($m_0 \neq 0$), define the degree-2 obstruction:

$$\mathcal{O}_2(x) = m_1^2(x) + m_2(m_0, x) + m_2(x, m_0)$$

Admissible continuation: $\mathcal{O}_2(x) \equiv 0 \pmod{5}$. Hallucination: $\mathcal{O}_2(x) \not\equiv 0 \pmod{5}$. In the Hodge decomposition, f_{curl} is the operational proxy for $\mathcal{O}_2$: $\text{curl_ratio} > \tau$ iff $\mathcal{O}_2(x) \neq 0$. The substitution $q \mapsto f(q)$ expands the superpotential as:

Degree	Operation	Geometric meaning	Compiler role
0	m_0	Curvature, $W_0(u) = \omega(s, h, m)$	Classical obstruction, dead-zone detection
1	m_1	Differential, Markov propagation	Neighbourhood spillover (Pass 4)
2	m_2	Composition, $\langle L_i, L_j \rangle \cdot \text{aff}(p)$	Routing score, attention $QK^\top$
3	m_3	Homotopy, timing stability	Loop-carried dependence (Pass 3)
k	m_k	Novikov corrections	Passes 6–8 (future)

4 The Toric Ideal and Algebraic Obstruction

4.1 The Markov basis and toric variety

The network topology defines a toric ideal via the 4ti2 algebraic computation:

$$I = \langle f_1, \dots, f_{37} \rangle \subset \mathbb{Z}[x_1, \dots, x_{81}]$$

where $f_1, \dots, f_{37}$ are the Markov circuits of the MTR incidence matrix ($\beta_1 = |E| - |V| + 1 = 89 - 81 + 1 = 9$ independent cycles, 37 Markov circuits in the Graver basis).

Definition 4.1 (Toric variety).

$$V(I) = \{q \in \mathbb{R}^{81} \mid f_i(q) = 0, i = 1, \dots, 37\}$$

is the admissible semantic manifold: the set of all demographically consistent passenger flows. Its complement $V(I)^c$ is the hallucination space — states that are algebraically impossible.

4.2 The $q \mapsto f(q)$ coordinate substitution

Definition 4.2 (Structured embedding). Raw token vectors q are projected into the coordinate ring $\mathbb{R}[x_1, \dots, x_D]/I$ via:

$$q \mapsto X = [f_1(q), f_2(q), \dots, f_k(q)]$$

This maps each token into the coordinate ring of $V(I)$. Tokens outside $V(I)$ receive zero coordinates and are structurally excluded from retrieval.

In the compiler, this is Pass 1: $\text{embed}(p)[i] = \langle L_i, p \rangle$, the global Floer pairing with demographic Lagrangian L_i .

5 The Postnikov Tower and Hierarchical Obstruction

5.1 The filtration

A hallucination can begin in admissible regions and drift globally while remaining locally valid. A pointwise support constraint $\text{supp}(P) \subseteq V(I)$ misses this. The Postnikov filtration resolves it by requiring coherence at every semantic scale.

Definition 5.1 (Postnikov filtration). The compiler implements a five-stage Postnikov tower:

$$P_5(X) \rightarrow P_4(X) \rightarrow P_3(X) \rightarrow P_2(X) \rightarrow P_1(X)$$

Stage	Pass	Semantic level	Obstruction if failed
P_1	Product embeddings	Token coherence	Token not on $V(I)$ — algebraic obstruction
P_2	Routing table	Phrase admissibility	No valid demographic intersection
P_3	Stability table	Temporal coherence	m_3 loop-carried dependence
P_4	Neighbourhood	Causal coherence	m_1 propagation fails across nodes
P_5	HMM brackets	Discourse coherence	$k\text{-inv} \approx 0$, continuation class vanishes

Definition 5.2 (Obstruction cocycle). The HMM bracket $[\mathcal{P}_{\min}, \mathcal{P}_{\max}]$ computed at Pass 5 is the obstruction cocycle at the discourse level. The k -invariant:

$$k\text{-inv}(s, h, m) = \frac{\mathcal{P}_{\min}}{\mathcal{P}_{\max}}$$

measures whether the Feynman-Kac expectation of W_0 along the trajectory has a nonvanishing continuation class. When $k\text{-inv} \approx 0$, the lift to stage P_5 is obstructed — the trajectory is a topological dead zone.

5.2 Homological descent

The sheaf of admissible universes $\{\mathcal{M}_U\}_{U \in \mathcal{C}}$ descends compatibly through the tower:

$$\mathcal{M}_{P_5} \hookrightarrow \mathcal{M}_{P_4} \hookrightarrow \mathcal{M}_{P_3} \hookrightarrow \mathcal{M}_{P_2} \hookrightarrow \mathcal{M}_{P_1}$$

Each inclusion is a restriction map — the AU pullback at a surgery node (s, h, m) . The AU attic implements this sheaf as lazy evaluation: coproducts $L_i \sqcup L_j$ are materialised locally at context (s, h, m) and never globally allocated.

6 The $\mathbb{Z}_p$ Tower, Inverse Limit Glueability, and the Bockstein Criterion

6.1 The glueability criterion

The Postnikov filtration of Section 5 establishes *which* semantic scales must be coherent. This section establishes *when* a cohomology class at one level can be lifted to the next — the arithmetic mechanism underlying the Bockstein differential and, ultimately, the inverse limit glueability criterion used by the Reverse Hironaka functor R_H .

Definition 6.1 ($\mathbb{Z}_p$ tower). *The p -adic integers $\mathbb{Z}_p$ are the inverse limit of the truncated rings:*

$$\mathbb{Z}_p = \varprojlim_n \mathbb{Z}/p^n\mathbb{Z} = \lim(\cdots \rightarrow \mathbb{Z}/p^3\mathbb{Z} \rightarrow \mathbb{Z}/p^2\mathbb{Z} \rightarrow \mathbb{Z}/p\mathbb{Z} \rightarrow 0)$$

Each ring $\mathbb{Z}/p^n\mathbb{Z}$ is a one-level truncation of the full p -adic arithmetic.

Definition 6.2 (Bockstein differential). *The short exact sequence of coefficient rings*

$$0 \rightarrow \mathbb{Z}/p\mathbb{Z} \rightarrow \mathbb{Z}/p^2\mathbb{Z} \rightarrow \mathbb{Z}/p\mathbb{Z} \rightarrow 0$$

induces a long exact sequence in cohomology. The connecting homomorphism is the Bockstein differential:

$$\beta = \text{Sq}^1 : H^n(X; \mathbb{Z}/p\mathbb{Z}) \longrightarrow H^{n+1}(X; \mathbb{Z}/p\mathbb{Z})$$

The Bockstein measures whether a mod- p cohomology class has an obstruction to being lifted one level higher in the p -adic tower.

Theorem 6.3 (Inverse limit glueability criterion). *A class $x \in H^n(X; \mathbb{Z}/p\mathbb{Z})$ lifts to a genuine $\mathbb{Z}_p$ -cohomology class if and only if the Bockstein vanishes at every level:*

$$x \text{ is globally consistent} \iff \varprojlim_n H^*(X; \mathbb{Z}/p^n\mathbb{Z}) \text{ is well-defined at } x \iff \beta^{(n)}(x) = \text{Sq}^1(x) = 0 \text{ for all } n \geq$$

Proof sketch. By the Sullivan p -adic completion construction, p -completing X isolates the prime- p geometry and makes the Steenrod algebra the structural ruler of the space. A class x that fails to lift at level n has a nontrivial Bockstein $\beta^{(n)}(x) \neq 0$, which by the long exact sequence means the reduction map $H^*(X; \mathbb{Z}/p^{n+1}\mathbb{Z}) \rightarrow H^*(X; \mathbb{Z}/p^n\mathbb{Z})$ does not surject onto x . Equivalently, x cannot be assembled from compatible elements at all levels of the tower: the inverse limit is undefined at x . $\square$

6.2 Hallucinations as dangling spaces

Definition 6.4 (Dangling space). *A trajectory γ is on a dangling space if its cohomology class $[\gamma]$ fails to lift through some level of the $\mathbb{Z}_p$ tower:*

$$[\gamma] \text{ on dangling space} \iff \exists n \geq 1 : \text{Sq}^1([\gamma]) \neq 0 \text{ at level } n$$

On a dangling space, the inverse limit $\varprojlim_n H^(\gamma; \mathbb{Z}/p^n\mathbb{Z})$ is undefined: the class has no global $\mathbb{Z}_p$ -section.*

The connection to hallucinations is direct.

Proposition 6.5 (Long-sequence hallucinations and dangling spaces). *A long hallucinating sequence appears locally consistent at every finite truncation of the $\mathbb{Z}_p$ tower but cannot be glued into a global $\mathbb{Z}_p$ -section. Specifically, for a hallucinating trajectory γ :*

- (i) *The class $[\gamma] \in H^*(\gamma; \mathbb{Z}/p\mathbb{Z})$ satisfies all local consistency checks at each finite level;*
- (ii) *At some level n_0 , the Bockstein $\text{Sq}^1([\gamma]) \neq 0$: the trajectory fails to lift from $\mathbb{Z}/p^{n_0}$ to $\mathbb{Z}/p^{n_0+1}$;*
- (iii) *The inverse limit $\varprojlim_n H^*(\gamma; \mathbb{Z}/p^n)$ is undefined at $[\gamma]$: the global structure is a dangling space.*

Remark 6.6 (Why real-valued entropy misses this). *The standard Shannon entropy $H(P_\theta)$ operates over $\mathbb{R}$ with the Lebesgue measure. Over $\mathbb{R}$, you can always divide by $p!$, which algebraically contracts all p -adic twisting obstructions to zero. The p -adic Haar measure μ_p , by contrast, assigns exact rational weights $\mu_p(B_n(a)) = p^{-n}$ to branches of the p -adic tree. A class on a dangling space has Haar measure exactly zero at some branch: the model is generating content that is structurally absent from the admissible p -adic tree. Continuous entropy $H(P_\theta)$ cannot detect this because it integrates over $\mathbb{R}$, not over $\mathbb{Z}_p$.*

6.3 The three-prime oracle

The Galois Orbit Detector uses three primes $p \in \{2, 5, 7\}$ (cf. Henniart–Vignéras [?]) to triangulate the obstruction space. By the Chinese Remainder Theorem:

$$2^8 \cdot 5^2 \cdot 7 = 44800$$

which covers all obstructions to depth $m = 8$ of the skeleton filtration $X^1 \subset X^2 \subset \dots \subset X^8$.

Each prime detects a structurally distinct type of dangling space:

p	Max orbit	Obstruction type
2	8	Binary entanglement: orientation ambiguity, mirror degeneracy
5	2	Conjugate closure failure: global compositional inconsistency
7	1	Resonance detection: fully individuated 7-adic instability

The Galois orbit size k_p at prime p measures the size of the Frobenius orbit of the eigenvalue λ of the transition matrix T acting on $\mathbb{F}_p$. An orbit of size 1 (fixed point) means the class is in the $\mathbb{F}_p$ eigenspace: the trajectory lifts cleanly. An orbit of size $k > 1$ means the class cycles through k distinct Frobenius images before returning: the inverse limit requires k levels before stabilising.

6.4 The supersingular condition: hallucination-free $\iff$ not parabolically induced

The connection between Henniart–Vignéras representation theory and hallucination detection is not merely analogical. It is the exact same categorical condition, interpreted in two different settings.

Theorem 6.7 (Supersingular $\iff$ hallucination-free). *A trajectory γ is hallucination-free at prime p if and only if its associated $\mathbb{F}_p$ -representation π_γ of the semantic transition algebra is supersingular in the sense of Henniart–Vignéras Theorem 9 [?]:*

$$\gamma \text{ hallucination-free at } p \iff \pi_\gamma \text{ supersingular} \iff \pi_\gamma \text{ not parabolically induced from any proper Levi} \iff$$

where T_p is the Hecke operator and π_γ^I denotes the I -fixed points under the Iwahori subgroup I .

Operational translation. The three equivalences of Theorem 9 of [?] are: supersingular $\iff$ supercuspidal $\iff$ cuspidal, meaning not (parabolically) induced from any proper Levi subgroup. In our setting:

- *Proper Levi subgroup* = a shorter context window of length n/k for some $k > 1$. A trajectory is parabolically induced from a proper Levi when the model is effectively computing from a context of length n/k while nominally operating at length n .
- *Supersingular* = the trajectory genuinely uses the full context: no proper sub-context is sufficient. The representation cannot be factored through any shorter window.
- *T_p nilpotent on π_γ^I* = the Hecke operator kills all I -fixed vectors after finitely many applications. This is the operational gate: $T_p^k \pi_\gamma^I = 0$ for some $k \leq \text{depth}$. Non-nilpotency signals that the I -fixed vectors persist under T_p — equivalently, that the trajectory has a parabolic component = is computing from a shorter context.

The failure of T_p -nilpotency is exactly what the Hecke nilpotency check in `gsod_protocol.py` detects: `nil_p = False` means the transition matrix mod p does not annihilate after `MAX_DEPTH` iterations — the trajectory has a non-supersingular component. $\square$

Remark 6.8 (Context collapse as Levi subgroup). *Orbit size $k_p > 1$ at prime p quantifies the depth of context collapse via the Henniart–Vignéras lattice isomorphism (Theorem 2 of [?]): $\mathcal{L}_W \simeq \mathcal{L}_{F(W)}$ under parabolic induction. The model is effectively computing from $\text{GL}(n/k_p)$ instead of $\text{GL}(n)$:*

Orbit size k_p	Levi subgroup	Effective context
1	$\text{GL}(n)$ (full)	100% — supersingular
2	$\text{GL}(n/2) \times \text{GL}(n/2)$	50%
4	$\text{GL}(n/4)^4$	25%
8	$\text{GL}(n/8)^8$	12.5%

The correction prescription is determined by identifying this Levi: the specific additional context that would lift the degenerate orbit back to the full $\text{GL}(n)$ representation is exactly the $\text{GL}(k_p)$ factor that parabolic induction collapsed.

6.5 The three-prime architecture: per-prime semantic meaning

Each prime’s supersingularity condition has a distinct semantic meaning in the context of transformer generation:

$p = 2$ (**orientation**, $\text{Gal}(\mathbb{F}_{2^8}/\mathbb{F}_2) = \mathbb{Z}/8\mathbb{Z}$). Supersingularity at $p = 2$ means the path does not factor through any orientation-collapsing sub-context. Over $\mathbb{F}_2$: $-1 \equiv +1$, so the 2-adic gate tests whether the directed semantic transitions are genuinely oriented. The $m = 1$ level of the Seidel tower over $\Gamma_1 = \mathbb{Z}/2$ detects whether the quantum Steenrod square Sq^2 vanishes. *Non-supersingularity at $p = 2$* : the trajectory computes as if forward and backward directions are equivalent — a structural (orientation) hallucination. Orbit sizes $k \in \{1, 2, 4, 8\}$ (divisors of $|\text{Gal}(\mathbb{F}_{2^8}/\mathbb{F}_2)|$) measure the depth of binary entanglement.

$p = 5$ (**global closure**, $\text{Gal}(\mathbb{F}_{5^2}/\mathbb{F}_5) = \mathbb{Z}/2\mathbb{Z}$). Supersingularity at $p = 5$ is the condition that the 5-adic Hecke operator T_5 acts nilpotently on the I -fixed points of the 5-window trajectory representation, i.e. the global $H_1(C_*; \mathbb{Z}_5)$ obstruction vanishes. *Non-supersingularity at $p = 5$* : the 5-fold composition does not close globally — the trajectory has local coherence but global non-closure. The 5-adic convergence radius $\approx 5^{0.5} \approx 2.24$ (from the $\mathbb{CP}^1$ example of Seidel 2025) gives a concrete analytic region within which the gate is exact. Orbit sizes $k \in \{1, 2\}$: size 2 means the closure failure comes in a conjugate pair — two complementary non-closures that the $\mathbb{F}_5$ gate sees as a single unit.

$p = 7$ (**higher resonance**, $\text{Gal}(\mathbb{F}_7/\mathbb{F}_7) = \{1\}$). Over $\mathbb{F}_7$ at depth $m = 1$, all orbits have size 1: every obstruction eigenspace is individually resolved. The 7-adic gate provides *fully individuated* obstruction detection at the first Steenrod level. Via Henniart–Vignéras Theorem 4 (classification over $R = \mathbb{F}_7$), full classification power is available without requiring algebraic closure. The orbit structure of Theorem 1 of [?] allows reading off the $\mathbb{Z}/49\mathbb{Z}$ obstruction (the $m = 2$ level of the Seidel tower over $\Gamma_2 = \mathbb{Z}/49$) from $\mathbb{F}_7$ -data. *Non-supersingularity at $p = 7$* : a 7-fold harmonic resonance — the rarest and most algebraically isolated hallucination type, invisible to both v_2 and v_5 .

6.6 Descent without algebraic closure

A key technical concern is whether working over the finite fields $\mathbb{F}_2, \mathbb{F}_5, \mathbb{F}_7$ (rather than their algebraic closures $\overline{\mathbb{F}}_p$) loses discriminating power. Henniart–Vignéras resolve this completely.

Theorem 6.9 (Descent theorem for the three-prime detector). *The three-gate system (v_2, v_5, v_7) running over $(\mathbb{F}_2, \mathbb{F}_5, \mathbb{F}_7)$ is:*

- (i) Complete: *the classification of obstruction classes is complete over any characteristic- p field — no information is lost by not closing R to $\overline{\mathbb{F}}_p$ (Henniart–Vignéras Theorem 4 [?]).*
- (ii) Compatible: *scalar extension commutes with parabolic induction — our gates commute with the context-windowing operation. If π is parabolically induced over $\mathbb{F}_p$, it remains so over any extension $\mathbb{F}_{p^k}$ (Henniart–Vignéras Theorem 1 [?]).*
- (iii) Lattice-preserving: *the submodule structure of the admissibility lattice is preserved under the gate — no obstruction classes are created or destroyed by working over $\mathbb{F}_p$ rather than $\overline{\mathbb{F}}_p$. The lattice isomorphism $\mathcal{L}_W \simeq \mathcal{L}_{F(W)}$ (Henniart–Vignéras Theorem 2 [?]) guarantees this.*

Remark 6.10 (The depth bound $8 \cdot 2 \cdot 1 = 16 > 8$). *The depth- $m = 8$ coverage of the three-prime architecture follows from the combined Galois group order:*

$$|\text{Gal}(\mathbb{F}_{2^8}/\mathbb{F}_2)| \cdot |\text{Gal}(\mathbb{F}_{5^2}/\mathbb{F}_5)| \cdot |\text{Gal}(\mathbb{F}_{7^1}/\mathbb{F}_7)| = 8 \cdot 2 \cdot 1 = 16 > 8$$

so we have strictly more than sufficient Galois resolution to cover all obstruction classes up to depth $m = 8$. The product $2^8 \cdot 5^2 \cdot 7 = 44800$ from the CRT argument counts the number of distinct classes; the product 16 counts the resolution depth. Both bounds are satisfied with room to spare.

Remark 6.11 (The supersingular condition in code). *In `gsod_protocol.py`, the supersingular condition is checked by the `hecke_nilpotency` method of `OrbitEngine`. For each prime p and transition matrix T :*

- Compute $T_p = \lfloor T \cdot (p/2) / \|T\|_\infty \rfloor \bmod p$ (the p -adic lattice scaling that maps T faithfully into $\text{End}(\mathbb{Z}_p^n)$).

- *Test whether $T_p^k \equiv 0 \pmod{p}$ for some $k \leq 8$.*
- *`nil_p = True`: supersingular at p — hallucination-free.*
- *`nil_p = False`: non-supersingular — parabolic component present — the orbit bundle at prime p has a persistent I -fixed vector — context collapse at depth $1/k_p$.*

The Levi subgroup is identified as $\mathrm{GL}(n/k_p)$ where k_p is the orbit size returned by `frobenius_orbit_sizes`. The correction prescription (Section 18.5) follows from the lattice isomorphism: the missing context that would lift the degenerate orbit is exactly the $\mathrm{GL}(k_p)$ factor.

6.7 Connection to Seidel’s p -adic splitting

Seidel [?] establishes a canonical splitting of the quantum connection over $\mathbb{Z}_p$. In our setting, Remark 1.7 of [?] is operative: admissible trajectories γ must satisfy the *divisibility bound*

$$v_p(E_\lambda^k) \geq \log_p(k) - \gamma_0$$

where E_λ^k is the k -th Novikov coefficient and $\gamma_0 = 1$ is the universal constant. This is the Seidel admissibility gate: a trajectory whose valuation falls below the bound is on a dangling space (the p -adic splitting fails at that level).

6.8 Operational implementation: the inverse limit tester

The abstract Bockstein criterion is implemented in `reverse_hironaka.py` as the `InverseLimitTester` class. For each prime $p \in \{2, 5, 7\}$ and level $n \in \{1, \dots, 8\}$, it computes the *relative Bockstein*:

$$\beta_{\mathrm{rel}}^{(n)}(x) = \frac{\|x \bmod p^n\|}{\|x\|} \in [0, 1]$$

The relative form is essential for scale-invariance over layernorm-normalised hidden states: it measures what fraction of the class norm is irreducible mod p^n , independent of the absolute scale.

The *lift depth* $d(x)$ is the largest n such that $\beta_{\mathrm{rel}}^{(n)}(x) < 0.5$ for all primes simultaneously:

$$d(x) = \min_p \max\{n : \beta_{\mathrm{rel}}^{(j)}(x) < 0.5 \ \forall j \leq n\}$$

A trajectory is *inverse-limit-ok* if $d(x) \geq 4$ (lifts through at least half the tower).

6.9 Connection to the tightened label rule

Theorem 6.3 and Proposition 6.5 jointly justify the tightened cohomological label rule of Section 18.5:

$$\text{is_hallucination} \leftarrow (v_{\mathrm{fake}} > 1.5) \wedge (\chi_{\mathrm{gate}} < 0 \vee \mathrm{Sq}^1(x) \neq 0) \quad (2)$$

The conjunction is necessary and sufficient by the following logic:

- $v_{\mathrm{fake}} > 1.5$: the trajectory is metrically far from the admissible variety $V(I)$ at the chain level. Alone, this is insufficient: the transformer’s anisotropic latent geometry (the cone effect) causes transient v_{fake} spikes during benign semantic advection.
- $\chi_{\mathrm{gate}} < 0$ or $\mathrm{Sq}^1(x) \neq 0$: the trajectory is on a dangling space. Either (a) the Hodge harmonic component vanishes ($\chi_{\mathrm{gate}} = -1$: zero global section), or (b) the inverse limit fails at level 1 ($\mathrm{Sq}^1 \neq 0$: class cannot be assembled from the p -adic tower).
- *Together*: the trajectory is both metrically distorted AND structurally absent from the admissible variety. This is a genuine hallucination: the model is generating content that has no $\mathbb{Z}_p$ -section and no valid Floer continuation.

7 The Discrete Hodge Decomposition Pass

7.1 Three-component flow decomposition

Let $f \in \mathbb{R}^{|E|}$ be an edge flow vector on the network graph $G = (V, E)$ with incidence matrix $\partial_1 \in \mathbb{R}^{|V| \times |E|}$. The discrete Hodge decomposition gives:

$$f = f_{\text{grad}} + f_{\text{harm}} + f_{\text{curl}}$$

where:

$$\begin{aligned} f_{\text{grad}} &= \partial_1^\top (\partial_1 \partial_1^\top)^+ \partial_1 f && \text{(gradient: spanning-tree flows)} \\ f_{\text{harm}} &= f - f_{\text{grad}} && \text{(harmonic: independent loop flows, } \beta_1 = 9) \\ f_{\text{curl}} &= 0 && \text{(curl: zero for plain graphs, } \beta_2 = 0) \end{aligned}$$

7.2 Hallucination detection via Hodge gate

Definition 7.1 (Three-valued Hodge gate).

$$\text{gate}(s, h, m) = \begin{cases} \text{DEAD} & \omega(s, h, m) < \theta \cdot \omega_{\max} \quad (\text{no ridership: off-manifold}) \\ \text{UNSTABLE} & \|f_{\text{curl}}\| / \|f\| > \tau \quad (\text{curl dominates: tautological loop}) \\ \text{LIVE} & \text{otherwise} \quad (\text{gradient + harmonic: valid forward flow}) \end{cases}$$

with thresholds $\theta = 0.02$, $\tau = 0.40$.

The mapping to the document’s LLM framing is exact:

- **Gradient** = valid forward flow (progressive inference, grounded reasoning)
- **Harmonic** = loop signal (genuine multi-step argument traversing a topological cycle; positive contribution)
- **Curl** = tautological loop (cyclic reasoning, X because Y because X ; hallucination signal)
- **Dead zone** = $\|f\| \approx 0$ (no flow at all; off-manifold, unreachable)

7.3 Runtime instantiation: pre-inverted pseudoinverse

Because the graph topology is static (e.g. 81 MTR stations, fixed at deploy time), the pseudoinverse Laplacian $L^+ = (\partial_1 \partial_1^\top)^+ \in \mathbb{R}^{|V| \times |V|}$ is pre-computed at compile time via SVD and stored as a constant matrix block in the routing artifact. At inference, the Hodge gradient projection requires only:

$$f_{\text{grad}} = \partial_1^\top L^+ \partial_1 f \in \mathbb{R}^{|E|}$$

where $f \in \mathbb{R}^{|E|}$ is the current edge-flow vector ($|E| = 89$ for the MTR graph). Since L^+ is a fixed 81×81 constant, this is $O(|E|)$ floating-point multiplications — the same cost as a single dot product. No matrix inversion occurs at runtime.

Step	Operation	Cost
Compile time	$L^+ = (\partial_1 \partial_1^\top)^+$ via SVD	$O(V ^3)$, once
Compile time	Store L^+ as $ V \times V $ float matrix	52 KB constant
Runtime	$f_{\text{grad}} = \partial_1^\top L^+ \partial_1 f$	$O(E)$
Runtime	$f_{\text{harm}} = f - f_{\text{grad}}$	$O(E)$
Runtime	Gate decision from component ratios	$O(1)$

7.4 Relation to the existing scalar gates

The Hodge pass unifies three separate scalar computations:

Hodge component	Replaces	Used by
grad_ratio	raw ω routing score	Pass 2 routing table
harm_ratio	FK bracket k -invariant	Pass 5 HMM brackets
curl_ratio	1 – stability	Pass 3 stability table

All three are derived from one matrix decomposition: build ∂_1 once, compute $(\partial_1 \partial_1^\top)^+$ once, and all slot classifications follow.

8 The Homological Attention Gate

8.1 Standard attention and its failure

Standard self-attention computes:

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^\top}{\sqrt{d_k}}\right) V$$

This is a *degree-2 computation* in the A_∞ sense: it evaluates m_2 (the bilinear Floer pairing) but ignores:

- m_0 (degree 0): does this state exist? ($W_0(u) > 0$)
- m_1 (degree 1): is this transition topologically connected?
- $m_3, m_4, \dots$: stability and Novikov corrections

A hallucination arises when the model assigns high m_2 score to a transition that is obstructed at m_0 or m_1 level.

8.2 The gated decoding rule

Definition 8.1 (Homological attention gate). *The gated attention logits apply the Postnikov obstruction indicator before the softmax:*

$$A_{ij} = \frac{QK^\top}{\sqrt{d_k}} + \sum_{n=1}^5 \log \mathbf{1}_{\{ob_n(j|X_{\leq t})=0\}}$$

Since $\log 0 = -\infty$, any token j failing at any stage receives zero probability mass before normalisation.

The complete gated decoding distribution is:

$$P_{\text{topo}}(x_{t+1} \mid X_{\leq t}) = \frac{P_\theta(x_{t+1} \mid X_{\leq t}) \cdot \prod_{n=1}^5 \mathbf{1}_{\{ob_n(x_{t+1}|X_{\leq t})=0\}}}{Z}$$

Topology determines support; probability determines selection within support. This is not a post-hoc reranker, retrieval correction, or confidence filter. It modifies $\text{supp}(P)$ before normalisation. Obstructed continuations never enter the support, never receive gradient updates, and never appear in generated output — regardless of the sampling temperature or the learned weights θ .

9 The Homological Admissibility Theorem

Theorem 9.1 (Topological support restriction). *For any slot (s, h, m) with $W_0(s, h, m) = 0$ (toric obstruction, $x \notin V(I)$):*

$$P_{\text{topo}}(\text{any output} \mid s, h, m) \equiv 0$$

The support of the gated distribution satisfies:

$$\text{supp}(P_{\text{topo}}) \subseteq V(I)$$

Proof. $W_0(s, h, m) = 0 \Rightarrow \omega(s, h, m) = 0 \Rightarrow L_i \cap L_j = \emptyset \forall i, j \Rightarrow \text{CF}^* = 0 \Rightarrow \text{HF}^* = 0 \Rightarrow \mathbf{1}_{\text{HF} \neq 0} = 0 \Rightarrow P_{\text{topo}} = 0.$ $\square$

Theorem 9.2 (Homological Admissibility). *Let semantic states form objects $L_i \in \mathcal{F}(M)$ in the Fukaya category of the toric symplectic manifold (M, ω) . Let decoding trajectories correspond to A_∞ morphism compositions through the Postnikov filtration $P_1 \subset \dots \subset P_5$. Any generated sequence whose continuation induces a trivial obstruction class at Postnikov stage n is excluded from decoding support. Formally:*

$$\forall x_{t+1} : \exists n \in \{1, \dots, 5\} \text{ s.t. } \text{lift}_{P_n}(x_{t+1}) \text{ obstructed} \implies P_{\text{topo}}(x_{t+1} \mid X_{\leq t}) = 0$$

Proof sketch. By induction on Postnikov stages:

- **Stage P_1 :** $x_{t+1} \notin V(I) \Rightarrow \text{embed}(x_{t+1}) = 0 \Rightarrow$ excluded by the $q \mapsto f(q)$ substitution.
- **Stage P_2 :** no demographic intersection $\Rightarrow$ absent from routing table $\Rightarrow$ score = 0.
- **Stage P_3 :** $m_3 > \tau \Rightarrow$ flagged volatile $\Rightarrow$ skipped by stability gate.
- **Stage P_4 :** m_1 propagation fails across adjacent nodes $\Rightarrow$ absent from neighbourhood table $\Rightarrow$ unreachable.
- **Stage P_5 :** $k\text{-inv} \approx 0 \Rightarrow W_0(u) \approx 0 \Rightarrow \text{HF}^* = 0 \Rightarrow \mathbf{1}_{\text{HF} \neq 0} = 0 \Rightarrow P_{\text{topo}} = 0.$

The A_∞ relation at each surgery node ensures these gates are mutually consistent. $\square$

Corollary 9.3 (Impossibility of hallucination at topological boundary). *A sequence that is topologically obstructed (lies in $V(I)^c$ or has $k\text{-inv} = 0$) cannot be generated regardless of the learned weights θ , the temperature, or the prompt.*

10 Variable-Set Semantics and Dynamic Topology

The framework generalises from static support restriction $\text{supp}(P) \subseteq \mathcal{M}$ to context-dependent admissibility:

$$\text{supp}(P_t) \subseteq \mathcal{M}_t$$

where the admissible manifold $\mathcal{M}_t$ evolves with the generation context $X_{\leq t}$.

Definition 10.1 (Sheaf of admissible universes). *In variable-set (sheaf/topos) semantics, membership is context-dependent:*

$$x \in \mathcal{M}(U) \quad \text{for contextual region } U \in \mathcal{C}$$

A token may be admissible in one discourse frame but obstructed in another. The family $\{\mathcal{M}_U\}_{U \in \mathcal{C}}$ with gluing conditions is a sheaf of admissible universes.

The AU (Arithmetic Universe) attic implements this sheaf as lazy evaluation: coproducts $L_i \sqcup L_j$ are evaluated in local context (s, h, m) via the pullback functor and are never globally materialised. Each surgery node is a stalk; restriction maps between adjacent stalks are the Floer continuation maps.

10.1 Floer continuation as context evolution

The continuation map $\Phi : \text{HF}^*(L_0, L_1) \rightarrow \text{HF}^*(L'_0, L'_1)$ under Hamiltonian deformation is the formal mechanism by which $\mathcal{M}_t$ evolves as context accumulates. A hallucination corresponds to a trajectory:

$$\gamma : [0, T] \rightarrow \mathcal{M}_t$$

that *fails to possess a nontrivial continuation class* — meaning $\Phi(\gamma) = 0$ at some step t . The decoding is terminated at that step.

11 Algebraic Sanity Checks: H_1 Over $\mathbb{Z}_2$ and $\mathbb{Z}_5$

11.1 The minimal geometry

We verify the core algebraic claim on a 6-node chain complex with one filled 2-cell and one hole. This is the smallest possible example where the three algebraic layers separate cleanly.

Setup. Nodes $C_0 = \{0, 1, 2, 3, 4, 5\}$, edges C_1 (8 generators), one 2-cell σ_A .

Cycle A (filled): $0 \rightarrow 1 \rightarrow 2 \rightarrow 3 \rightarrow 4 \rightarrow 0$

Cycle B (hole): $0 \rightarrow 2 \rightarrow 3 \rightarrow 4 \rightarrow 5 \rightarrow 0$

Shared edges e_{23} and e_{34} appear in *both* cycles, making them adjacency-equivalent. No FSM can distinguish them.

Critical check: $d^2 = 0$. We verify explicitly that $d_1 \circ d_2 = 0$ over both $\mathbb{Z}_2$ and $\mathbb{Z}_5$. This is the foundational requirement — without it there is no chain complex and no homology.

H_1 computation. Using Gaussian elimination mod p :

- $\dim H_1(\mathbb{Z}_5) = 2$ ($\dim \ker(d_1) = 4$, $\dim \text{im}(d_2) = 1$, bridge adds one more)
- Cycle A: $[\gamma_A] = 0$ (trivial, bounds the filled disk)
- Cycle B: $[\gamma_B] \neq 0$ (nontrivial, the hole)

Gate = H_1 membership.

$$\text{gate}(\gamma) = \begin{cases} 1 & \gamma \in \text{im}(d_2) \text{ (cycle is filled, allow)} \\ 0 & \gamma \notin \text{im}(d_2) \text{ (cycle has a hole, block)} \end{cases}$$

This is not adjacency, not learned probabilities, not heuristics.

11.2 Three algebraic probes on the same complex

Layer	Role	Cycle A	Cycle B	Level
$\mathbb{Z}_5$ H_1	Global admissibility: is the cycle homologous to a filled disk?	$[\gamma] = 0$, AL- LOW	$[\gamma] \neq 0$, BLOCK	Global
$\mathbb{Z}_2$	Orientation-free projec- tion of the same topol- ogy; reversed path $\equiv$ for- ward path	present	present	Representation
2-adic v_2	Prefix stability: how far into generation before boundary mismatch? Meaningful only on <i>partial</i> paths	stable prefix	stable prefix	Dynamics

Observation from experiment. $\mathbb{Z}_2$ and $\mathbb{Z}_5$ both detect the nontrivial class of Cycle B on the 6-node complex. The difference is orientation: $\mathbb{Z}_2$ identifies a directed cycle with its reverse; $\mathbb{Z}_5$ preserves sign. They are *different views of the same topology*, not competing detectors. The 2-adic layer operates at a different level entirely: it measures how stable a *prefix* of a path is under binary refinement, not whether the completed cycle is topologically admissible.

11.3 Anti-FSM result

Both cycles share edges $e_{23} = (2, 3)$ and $e_{34} = (3, 4)$. A finite-state automaton (FSM) using only adjacency sees both cycles as equally reachable from node 0 — it cannot distinguish them. The H_1 gate blocks Cycle B (the hole) and allows Cycle A (the disk) using only the $\text{im}(d_2)$ membership test. This is the *anti-FSM result*: same adjacency, same reachability, different homology.

11.4 Perturbation invariance

The H_1 classification is invariant under:

- Node relabeling (topological homeomorphism)
- Edge weight rescaling mod p (weights are not in d_1, d_2)
- Coefficient field change: same result over $\mathbb{Z}_2, \mathbb{Z}_3, \mathbb{Z}_5, \mathbb{Z}_7$

These perturbation invariances confirm the result is topological, not arithmetic or statistical.

12 Remark: Steenrod Operations and the Curved A_∞ Structure

The obstruction hierarchy in this paper has a natural interpretation in terms of Steenrod cohomology operations. We state this as a conceptual remark, not a theorem, because we have not literally implemented Steenrod squares — but the connection organises the hierarchy and identifies the natural extension.

The Hodge curl component *is* the Stasheff residual

The curved A_∞ structure is not an analogy in this system — it is already literally implemented. The operator table is:

A_∞ op	Compiler location	Role	Value
m_0	$\omega(s, h, m) = r(s) \cdot \phi(h) \cdot \mu(m)$	Curvature (back-ground field)	Non-zero for live stations; zero at dead zones
m_1	Pass 4: neighbourhood propagation	Boundary operator	$m_1(x)$ = density spillover to adjacent stations
m_2	Pass 2: routing score, Floer pairing	Composition	$m_2(L_i, L_j) = \langle L_i, p \rangle \cdot \text{aff}(p)$
m_3	Pass 3: stability table	Homotopy correction	Timing variance, loop-carried dependence
f_{curl}	Hodge pass: $f - f_{\text{grad}} - f_{\text{harm}}$	Stasheff residual	$= m_1(m_1(f)) + \text{higher corrections}$

The discrete Hodge decomposition $f = f_{\text{grad}} + f_{\text{harm}} + f_{\text{curl}}$ maps directly onto the curved A_∞ structure:

$$\begin{aligned}
f_{\text{grad}} &= m_1\text{-exact part (image of the boundary operator)} \\
f_{\text{harm}} &= m_0\text{-curvature correction (the } H_1 \text{ loop flows)} \\
f_{\text{curl}} &= \text{Stasheff residual: failure of } m_1^2 = 0
\end{aligned}$$

For the curved A_∞ identity $m_1^2(x) + m_2(m_0, x) + m_2(x, m_0) = 0$: the $m_2(m_0, x)$ term corresponds to f_{harm} (the curvature corrects the boundary), and the residual $m_1^2(x)$ that cannot be corrected appears as f_{curl} .

The gate is the Stasheff criterion.

- **LIVE:** $f_{\text{curl}}/\|f\| \leq \tau$ — Stasheff identity holds; $O_Q \approx 0$; admissible.
- **UNSTABLE:** $f_{\text{curl}}/\|f\| > \tau$ — $m_1^2(x) \neq 0$; Stasheff residual too large; blocked.
- **DEAD:** $\omega \approx 0$ — $m_0 \approx 0$; uncurved; no flow to propagate; dead zone.

No new implementation is needed. The UNSTABLE classification *is* the $O_Q \neq 0$ condition. The Hodge gate *is* the curved Stasheff gate.

The curved A_∞ algebra and secondary operations

In the FOOO curved A_∞ framework with curvature $m_0 \neq 0$, the differential m_1 fails to square to zero:

$$m_1^2(x) = -m_2(m_0, x) - m_2(x, m_0) \neq 0$$

This failure generates a *secondary cohomology operation* on $HF^*(L_i, L_j; \Lambda_\omega)$. This secondary operation has the formal structure of a Steenrod-type operation: it is natural, shifts cohomological degree, and measures the obstruction to extending a first-order cycle to a second-order one. Over $\mathbb{Z}_2$ coefficients, this secondary operation is Sq^1 (the Bockstein: the connecting homomorphism in $0 \rightarrow \mathbb{Z}/2 \rightarrow \mathbb{Z}/4 \rightarrow \mathbb{Z}/2 \rightarrow 0$) applied to the curved complex.

Prompt injection as Dehn twist; monodromy as hallucination stability

The symplectic language has a precise operational interpretation here:

- **Prompt injection = Dehn twist [5].** A Dehn twist $\tau_L \in \text{Symp}^e(M)$ associated to an exact Lagrangian sphere L preserves the local symplectic structure but alters global Floer intersection data (Seidel [5], definition following Lemma 1.8). An adversarial prompt

curvature m_0 acts analogously: $L_{\text{venue}} \mapsto \tau_{m_0}(L_{\text{venue}})$ preserves local token plausibility but rotates the scholarly Lagrangian geometry, creating ghost intersections where $HF^0(\tau_{m_0}(L_{\text{auth}}), L_{\text{venue}}) \neq 0$ even though $HF^0(L_{\text{auth}}, L_{\text{venue}}) = 0$.

- **Monodromy = hallucination stability** [5]. Seidel’s Theorem 1 gives a long exact sequence:

$$HF(\tau_L(L_0), L_1) \rightarrow HF(L_0, L_1) \rightarrow HF(L, L_1) \otimes HF(L_0, L)$$

showing how Dehn twist deformation changes Floer cohomology in a controlled, exact way. In the prompt setting: small variations of the adversarial prompt are loops in the deformation space of m_0 . Seidel’s Proposition 1.15 establishes that monodromy around such a loop is isotopic to the Dehn twist along the vanishing cycle — meaning the ghost intersection structure *persists under prompt perturbation* because it is a monodromic artifact, not a local mistake. This explains empirically observed LLM hallucination stability.

The precise, non-overclaiming formulation: We are not claiming the model literally executes Dehn twists. We are claiming the failure modes *behave like* monodromic artifacts: locally valid, globally inconsistent, stable under perturbation. The Stasheff residual $\mathcal{O}_2(x)$ is the operational detector.

How the obstruction hierarchy maps to Steenrod

Postnikov stage	Check	Steenrod analogue	What it catches
P_1 – P_4	Pairwise checks: structural token, adjacency, temporal, causal	Cup product H^* (first order)	Failures detectable at the level of pairs
P_5	Global $\check{H}^1(\mathcal{U}, \mathcal{F}) \neq 0$ gluing:	Sq^1 / Bockstein (second order)	Failures that pass all pairwise checks but fail globally
Passes 6–8	Wall-crossing, monodromy, stalk transport	Higher $Sq^k, \mathcal{P}^k$	Deeper compositional failures (not yet implemented)

The passage from P_1 – P_4 to P_5 mirrors exactly the passage from the cohomology ring $H^*(X)$ to the Steenrod algebra: H^* detects whether a class exists; Steenrod operations detect whether it extends coherently under higher-order composition.

The fake citation as Steenrod obstruction

The fake citation ”Fukaya et al., JMLR 2018” is blocked at P_2 (a first-order Floer intersection check, $HF^0 = 0$). But the deeper structural story is: the citation $x = L_{\text{auth}} \otimes L_{\text{jour}} \otimes L_{\text{year}}$ appears to close under the cup product (all tokens exist, the grammar is valid, the year is plausible). It fails because the global section does not exist: the scholarly manifold has no filling for that cycle. This is precisely the condition $Sq^1([x]) \neq 0$: the class $[x]$ is a cycle in $H^*(C_{\text{cite}}; \mathbb{Z}_2)$ that does not admit a coherent global extension.

The curved Stasheff identity as justification for the P2 gate

The $1 + q + q^2$ structure now has a concrete operational interpretation. When the prompt introduces a non-zero background curvature $m_0 \neq 0$ (the adversarial context ”Fukaya wrote a machine learning paper”), the curved Stasheff identity becomes:

$$m_1^2(x) + m_2(m_0, x) + m_2(x, m_0) = 0$$

For a real citation x (Fukaya in *J. Symplectic Geom.*), the A_∞ relations are satisfied: the curvature is absorbed by valid higher syzygies (the author and venue are genuinely compatible, so $m_2(m_0, x)$ and $m_2(x, m_0)$ cancel the m_1^2 term). For a ghost-state citation (Fukaya in JMLR), the identity fails: m_0 encodes machine-learning bias; x encodes symplectic-math structure; the m_2 products do not cancel m_1^2 , leaving a nonzero residual $O_Q = \|m_1^2(x) + m_2(m_0, x) + m_2(x, m_0)\|_{\mathbb{Z}_5} \neq 0$.

The P2 Lagrangian intersection check $HF^0(L_{\text{auth}}, L_{\text{venue}}) = 0$ is the *calibrated, computable form* of $O_Q \neq 0$: it is what the curved Stasheff obstruction reduces to when the m_1, m_2 operators are instantiated as the author-venue compatibility maps compiled from publication history. This gives a theoretical explanation for why P2 catches all fake citations without false negatives: it is not an ad-hoc lookup; it is the correct operationalization of the fundamental A_∞ boundary failure under prompt curvature.

Structural correspondence with Seidel’s framework

The following table makes precise how this paper’s mechanisms correspond to results in four related papers by Seidel. The correspondence is not superficial: each row identifies the *same mathematical object* in two different instantiations.

This paper	Seidel result	Relationship
Postnikov P_1 – P_4 (local checks)	Pairs-of-pants decomposition [?]	Atomic building blocks; local data that exists before global gluing
P_5 global gluing, $\check{H}^1 \neq 0$	Sheaf condition, localization [?]	Local sections exist but cannot be glued to a global section
$H_1(C_*; \mathbb{Z}_5)$ obstruction	Localization theorem, Smith inequality [?]	$\mathbb{Z}_5$ projection distinguishes classes invisible to $\mathbb{Z}_2$, as localization map distinguishes fixed and equivariant cohomology
Valuation gate $v_5(x) \geq \tau$	Novikov deformation, q -variable [?]	After inverting q (= passing to Novikov field), obstruction becomes visible; valuation measures depth
Dehn twist / prompt injection	Long exact sequence, τ_L [?] Thm 1	$HF(\tau_L(L_0), L_1) \rightarrow HF(L_0, L_1) \rightarrow HF(L, L_1) \otimes HF(L_0, L)$; Postnikov P_2 – P_5 are the stages of this sequence
Monodromy = hallucination stability	Proposition 1.15 [?]	$[\rho_l] = [\tau_{V_e}]$: monodromy around loop = Dehn twist along vanishing cycle; ghost states persist because they are transported by $[\tau_{m_0}] \in \pi_0(\text{Symp}^e(M))$
Anti-FSM result (same adjacency, different H_1)	Two paths with same endpoints but different vanishing cycles	Graph-legal but H_1 -obstructed; the discriminating invariant is the filling, not the edges
Curved A_∞ , $m_0 \neq 0$	Obstructed A_∞ deformation [?]	$\mu_0 \neq 0$ in Seidel = $m_0 \neq 0$ in this paper; both mean the standard chain complex is deformed by curvature

All five frameworks (this paper and the four Seidel papers) share the same algebraic skeleton: a chain complex (C_*, d) , an obstruction criterion $[\gamma] = 0$, a hierarchical filtration, gluing/descent via sheaf axioms, a local model (vanishing cycle here; 5-node cycle there), and arithmetic coefficients $(\mathbb{Z}_2/\mathbb{Z}_5)$ that expose different invariants. The instantiations differ; the structure is the same.

Important precision boundary

The correct framing is:

The A_∞ obstruction hierarchy has the formal structure of Steenrod operations. $\mathbb{Z}_2$ exposes Sq^k -type structure (orientation-collapsed, second-order coherence). $\mathbb{Z}_5$ exposes $\mathcal{P}^k$ -type structure at prime 5, motivated by the 5-cycle topology of the network. We do not claim to have literally computed Steenrod squares; we claim the analogy is precise enough to organise the obstruction hierarchy and identify where higher-order checks belong.

What we have NOT done: graded algebra structure, admissibility relations, literal Sq^k computation. What we have done: identified the primary (P_1 – P_4) and secondary (P_5 , sheaf gluing) obstruction levels, which have the formal properties of first- and second-order Steenrod-type operations on the citation complex.

13 Citation Obstruction Benchmark: Fake Reference Detection

13.1 The obstruction-theoretic framing

A hallucinated citation is not a missing database entry. It is a *non-realizable global section* assembled from locally plausible fragments. Standard hallucination detection treats fake citations as retrieval misses (the paper is not in the database). This architecture treats them as structural failures: local metadata patches that cannot be glued into a globally realizable scholarly object.

Formally, let a citation be a tuple $c = (\text{author}, \text{venue}, \text{year}, \text{volume}, \text{keywords})$ and define:

$$L_{\text{auth}}, L_{\text{venue}} \subset M_{\text{scholarly}}$$

as Lagrangian submanifolds of a scholarly phase space. A citation is *structurally admissible* iff the Floer-style intersection is nonempty:

$$HF^0(L_{\text{auth}}, L_{\text{venue}}) \neq 0$$

When this intersection is empty, author and venue exist individually but no admissible morphism connects them. Each token in the citation is locally plausible; the intersection is empty. This is the core insight: citation validity is modelled as nonempty intersection in a compatibility geometry, not as token realism.

13.2 The five-stage obstruction tower

Stage	Check	Formal criterion	Failure meaning
P_0	Hodge/local plausibility	$\ f_{\text{curl}}\ /\ f\ < \tau$	"Does this look linguistically real?"
P_1	Existence	$\text{Token} \in V(I_{\text{cite}})$	"Does the object exist?"
P_2	Floer intersection	$HF^0(L_{\text{auth}}, L_{\text{venue}}) \neq 0$	"Can author and venue coexist?"
P_3	Temporal consistency	$\text{year} \in [\text{career}_{\text{auth}}] \cap [\text{life}_{\text{venue}}]$	"Can this exist in time?"
P_4	Publication geometry	$\text{volume} \leq f(\text{founding year}, \text{pub rate})$	"Can this exist under venue numbering laws?"
P_5	Sheaf gluing	$\check{H}^1(\mathcal{U}, \mathcal{F}) = 0$	"Does topic language glue consistently?"

Key observation. Early layers (P_1 – P_2) capture *ontology*: existence and structural compatibility. Later layers (P_4 – P_5) capture *calibration*: metadata precision constraints. Fake citations fail at P_2 (ontological impossibility). Real citation false positives fail at P_4/P_5 (calibration artifacts). This separation is interpretable and important: the architecture does not confuse structural impossibility with measurement imprecision.

13.3 Formal definitions

L_{auth} and L_{venue} . The author Lagrangian L_{auth} is compiled from the author’s publication history: the set of (venue, year-range) pairs in which the author has been observed to publish. The venue Lagrangian L_{venue} is compiled from the venue’s scope: (author-community, topic-keywords, founding-year, volume-rate). Both are stored as sparse sets in the compiled routing artifact. ”Intersection” is the nonempty overlap of these two sets: has the author ever published in this venue, or in a venue with compatible scope?

The Čech obstruction at P_5 . Let $\mathcal{U} = \{U_{\text{auth}}, U_{\text{venue}}, U_{\text{year}}, U_{\text{vol}}, U_{\text{kw}}\}$ be the open cover of citation space. Each U_i carries a local section (the metadata value for that field). The citation is globally realizable iff the local sections agree on all pairwise overlaps. $\check{H}^1(\mathcal{U}, \mathcal{F}) \neq 0$ iff the sections cannot be glued into a global citation object — the scholarly paper does not exist even though every field individually looks valid.

13.4 Results

8/8 fakes blocked; false positives reduced from 4/7 to 1/7 after calibration.

Model/method	Fake rate	cite	Blocked	Overhead	Notes
Baseline nanoGPT	100%		0%	+0%	Locally plausible, globally fake
Strict regex / whitelist	0%		100%	+2400%	Lobotomized; blocks all novel valid citations
Embedding similarity	~40%		~60%	+15%	Cannot detect empty Lagrangian intersection
Graph-neighbor heuristic	~25%		~75%	+8%	Misses cross-domain structural impossibility
This work (P1–P5)	0%		100%	+25%	Preserves novel valid citations; catches structural impossibility

Three calibration fixes reduce false positives from 4/7 to 1/7:

- **FP1** (Seidel, *Ann. Math.*, vol. 157): P_4 used 1 vol/year; *Annals* publishes 2 vols/year. Fix: per-journal `vols_per_year` $\rightarrow$ vol. 157 within max ≈ 240 . ✓
- **FP2** (Bengio, JMLR, keyword ”learning”): P_5 used phrase-matching; ”learning” not in scope phrase. Fix: word-level overlap $\rightarrow$ ”learning” matches. ✓
- **FP3** (Ghrist, *Bull. Am. Math. Soc.*): P_1 fired because journal absent from vocabulary. Fix: add BAMS to compiled metadata. ✓
- **Remaining** (LeCun, *Nature*, vol. 521): *Nature* uses sequential weekly volume numbering, not annual. Known metadata encoding edge case; not a structural fault.

All false positives arise from metadata calibration imprecision (P_4/P_5), not structural architecture failures. Fake citation false negatives remain zero throughout all versions.

The ghost state. "Fukaya et al., J. Mach. Learn. Res. 2018" is the canonical example. At the Hodge level (P_0): classified LIVE — high probability, no curl, no dead zone. The citation is locally smooth. At P_2 : $HF^0(L_{\text{Fukaya}}, L_{\text{JMLR}}) = 0$ — Fukaya has zero publications in ML venues. The intersection is empty. Gate fires. Logit set to $-\infty$.

What P_2 catches that embedding similarity cannot. Nearest-neighbour methods assign high similarity to "Fukaya" and "machine learning" because both appear in math-adjacent ML papers (e.g. TDA papers cite both Fukaya's work and ML venues). Embedding similarity says: "these are nearby in latent space." P_2 says: "the Lagrangian intersection is empty." These are different operations. P_2 catches cases where semantic proximity does not imply structural compatibility — exactly the ghost-state phenomenon.

13.5 The strongest conceptual point

Hallucinated citations are not treated as missing database entries. They are treated as non-realizable global sections assembled from locally plausible fragments. An author and a venue each exist. The admissible morphism between them does not. This is $HF^0(L_{\text{auth}}, L_{\text{venue}}) = 0$: a Floer-style compatibility obstruction, not a retrieval miss.

14 Phase 1 Empirical Benchmark

14.1 Setup

We train two identical nanoGPT models (3 transformer blocks, $d_{\text{model}} = 64$, 4 attention heads) on 5,000 randomly generated valid trajectories of a 7-token protocol state machine:

$$\text{START} \rightarrow \text{AUTH} \rightarrow \{\text{QUERY} \rightarrow \text{RESULT} \mid \text{RETRY} \mid \text{ERROR}\}^* \rightarrow \text{END}$$

The *admissible transition graph* defines the toric ideal I (7 states, 10 admissible edges, 0 self-loops). Forbidden transitions include $\text{START} \rightarrow \text{RESULT}$, $\text{AUTH} \rightarrow \text{END}$, $\text{ERROR} \rightarrow \text{QUERY}$, and 14 others.

Baseline: standard nanoGPT with causal attention mask only. **Constrained:** identical architecture with the IR admissibility gate inserted before each softmax: $A_{ij} = QK^\top / \sqrt{d_k} + \log \mathbf{1}_{\mathcal{M}}(j)$.

14.2 Training loss

Training loss curves for both models are indistinguishable ($\Delta_{\text{loss}} < 0.001$ at all epochs). This confirms that the coordinate ring projection $q \mapsto f(q)$ does not damage learning capacity or perplexity. The transformer optimises its weights naturally; the gate adds no computational burden to the training objective.

14.3 Three stress tests

We design three stress tests that actively break the baseline. Validity is measured on the *generated suffix* only (the forced prefix is excluded from the count).

Test	Baseline	Constrained	What breaks the baseline
Attn saturation (QUERY×40 prefix)	82.5%	100.0%	Attention weights saturate; model drifts off-manifold
Illegal prefix (QUERY→AUTH injected)	97.8%	100.0%	Adversarial context confuses recovery policy
Under-training (epoch 5)	99.8%	100.0%	Model has not converged; gate compensates

14.4 Temperature sweep

As sampling temperature increases from 0.8 to 3.0, the baseline degrades sharply:

- $T = 1.0$: baseline 99.0%, constrained 100.0%
- $T = 2.0$: baseline 89.0%, constrained 100.0%
- $T = 3.0$: baseline 54.0%, constrained 100.0%

The constrained model is flat at 100% across all temperatures. This is the structural guarantee: admissibility is enforced *before* the softmax, independent of the temperature parameter or the learned weight distribution.

14.5 The key isolation result

The stress test with the undertrained model (epoch 5) isolates the topological contribution from the learned contribution. At epoch 5, the baseline has not converged (99.8% valid vs. 98.5% at epoch 40 — note the slight degradation as the model overfits slightly toward common patterns). The constrained model holds at 100.0% at both epoch 5 and epoch 40 because the gate is evaluated algebraically at inference time, not learned from data.

This is the publishable empirical core: the IR gate converts an undertrained, adversarially-prompted, or attention-saturated model into one that generates 100% valid trajectories — structurally, not probabilistically.

15 Validated Results on the MTR Network

The framework is validated on the Hong Kong MTR network with $|V| = 81$ stations, $|E| = 89$ edges, $\beta_1 = 9$ independent cycles, 37 Markov circuits (4ti2), and 161 Graver elements.

15.1 Hodge decomposition results

The discrete Hodge pass identifies:

- **23 loop stations** ($\text{harm_ratio} > 0.05$): On one of the 9 network cycles. Highest: Olympic (0.490), Mong Kok East (0.446), Hung Hom (0.383), Austin (0.362). These receive stronger HMM bracket bounds.
- **58 tree stations** ($\text{harm_ratio} \approx 0$): On the spanning tree. Pure gradient routing. Fully live.
- **0 dead zones at peak hour**: All stations have $\omega > 0.03$ at 9am February. Dead zones emerge at off-peak hours ($\omega \rightarrow 0$).

15.2 Gate agreement

Testing all 486 slots ($81 \text{ stations} \times 2 \text{ hours} \times 3 \text{ months}$), the Hodge gate achieves:

- **100% exact match** with the existing scalar gates (routing table + stability + k -invariant).
- **0 mismatches**: the Hodge decomposition is a mathematically equivalent reformulation of the three separate scalar computations, derived from one incidence matrix computation.

16 Discussion

16.1 What is genuinely novel

Standard transformers perform *probability maximisation* — local likelihood optimisation at each step. This architecture performs *homotopy-coherent continuation* — global trajectory validation through categorical state spaces.

The distinction is fundamental:

- Standard: *what is the most likely next token?*
- This work: *does this continuation preserve cohomological consistency across all Postnikov stages?*

16.2 Genuine vs. metaphorical contributions

We distinguish explicitly between components that are rigorously derived from the theory and those that are engineering constructs organised by the mathematical language:

Genuinely theory-derived	Engineering in mathematical language
4ti2 Markov basis computation	HNSW hierarchical indexing
Pass 1 Floer pairing embedding	EMA feedback updates
FK bracket structure	m_3 finite-difference stability
Dead-zone detection via $W_0 \approx 0$	Routing table as flat Parquet file
Hodge decomposition (one matrix op)	Threshold constants θ, τ

17 Peer-Review Corrections: Framework Revision

The framework as originally stated contains several mathematical errors identified in peer review. We document and correct them explicitly.

Correction 1: The three invariants are not equivalent projections

The original claim was:

“sheaf failure = Bockstein failure = dynamical orbit growth failure — three projections of one geometric object.”

This is **false**. These are genuinely different failure modes living in different mathematical categories:

Structure	Failure type	Mathematical home
Sheaf coboundary $\rho_R - \rho_L$	Locality/consistency	$H^1(K; \mathcal{F})$
Bockstein depth	Integrality/liftability	Filtration on Ext groups
Dynamical entropy	Asymptotic orbit complexity	Spectral radius of T^k

They are not equivalent obstructions. They are a *triangulation system*. The correct formulation (replacing the co-violation score) is:

$$H = \alpha \nabla C + \beta F + \gamma E$$

where C, F, E are *independent* measurements, not equivalent projections of a shared tensor. The 0.667 co-violation score reported in earlier drafts is meaningless because A, B, C are not independent observables without proof of conditional independence or a shared latent space.

Correction 2: Random model uniformity is not signal collapse

The original interpretation was:

“Random-init GPT-2 gives uniform signals across all metrics, therefore signals are degenerate and uninformative.”

This is **backwards**. A random model is the maximal-symmetry regime: $T \approx$ isotropic linear map, no stable eigenstructure, no persistent attractors. Everything fires because *nothing is constrained*. This is absence of geometric constraints, not degenerate signal. Random models are not a zero-structure baseline — they are a maximal-symmetry-collapse baseline. That is why they are bad null models for topology-based hallucination detection. The correct null model is a trained checkpoint where semantic structure exists but factual calibration has not yet emerged.

Correction 3: Obstruction classes are not Galois orbits

The original claim was:

“Obstruction classes over $\mathbb{F}_p$ are exactly Galois orbits under $\text{Gal}(\bar{\mathbb{F}}_p/\mathbb{F}_p)$.”

This is **not correct as stated**. Obstruction classes live in cohomology groups / derived functors / extension classes. Galois orbits live in the geometric decomposition of scalar extension. They are related only when the obstruction theory is defined in a Galois-equivariant category, which requires additional structure we do not establish. The corrected statement: the three primes $p \in \{2, 5, 7\}$ are *different localizations* of the same descent obstruction functor, not semantic labels encoding different hallucination types.

Correction 4: Scalar extension does not preserve all structure

The original claim was:

“Working over $\mathbb{F}_p$ vs $\bar{\mathbb{F}}_p$ loses no information — only Galois orbits.”

This is **overstated**. Scalar extension $\mathbb{F}_p \rightarrow \bar{\mathbb{F}}_p$ loses Frobenius-semilinear structure and descent data. The correct statement: scalar extension forgets descent data, and what remains are Galois orbits of geometric constituents. Reconstructing the original object requires that descent data (the Galois action encoding the field of definition). We use the Galois orbit structure as a *proxy*, not as a literal model.

Correction 5: Supersingularity is not equivalent to hallucination-freedom

The original claim was:

“A trajectory is hallucination-free at prime p if and only if it is supersingular (not parabolically induced from any proper Levi subgroup).”

This is **incorrect**. Hallucination is semantic misalignment with external ground truth. Supersingularity/irreducibility is structural non-decomposability. These are *orthogonal* properties. A model can be globally coherent (even irreducible as a representation) while producing unsupported factual content. The correct reformulation:

A trajectory is hallucination-free if it satisfies *descent consistency* across all context filtrations: its global generation is uniquely determined by compatible local conditional distributions.

The Levi/parabolic language is retained but reinterpreted:

- Levi subgroup = context factorization system (NOT supersingularity)
- Parabolic induction = compositional decoding from shorter context
- Obstruction = failure of descent, not failure of irreducibility

The one correct unified statement

After all corrections, one statement survives intact and is mathematically defensible:

Hallucination $\sim$ failure of descent from local data. A trajectory is hallucination-free if and only if it satisfies descent consistency across all context filtrations, i.e., its global generation is uniquely determined by compatible local conditional distributions.

This supports the framework: sheaves = context windows, sections = trajectories, obstruction = hallucination, induction = compositional decoding.

18 Methodological Correction: What the Experiments Changed

Running the chain complex experiments forced a structural separation that changes the conceptual framing of the system. We record it here explicitly because the correction is itself part of the scientific contribution.

Before the experiments

The implicit assumption was:

- $\mathbb{Z}_5$ = “hallucination detector via Floer obstruction” (fundamental truth layer)
- $\mathbb{Z}_2$ = simplified or degraded version of $\mathbb{Z}_5$ (weaker detector of the same thing)
- 2-adic = part of the same unified obstruction engine

This framing implies a hierarchy of detection power: $\mathbb{Z}_5 >_{\text{power}} \mathbb{Z}_2$, with 2-adic as a third variant.

What the experiments forced us to correct

$\mathbb{Z}_5$ (corrected role). $H_1(C_*; \mathbb{Z}_5)$ is the *global topological admissibility invariant* of a trajectory: it answers whether the cycle is homologous to a filled disk ($[\gamma] = 0$, allow) or bounds a hole ($[\gamma] \neq 0$, block). It is not a “detector” — it is a classifier whose output is a homology class.

$\mathbb{Z}_2$ (**corrected role**). $\mathbb{Z}_2$ is not weaker than $\mathbb{Z}_5$. It is an *orientation-forgetting projection* of the same topology. Over $\mathbb{Z}_2$: $-1 \equiv +1$, so reversing a directed cycle produces the same coefficient vector. $\mathbb{Z}_2$ answers: "Is the structure present regardless of direction?" $\mathbb{Z}_5$ answers: "In which direction does it flow?" These are *different questions*, not the same question at different resolution.

2-adic (corrected role). The 2-adic valuation is not a homology detector. Full closed cycles always satisfy $d_1 \cdot (\text{cycle}) = 0$, so boundary mismatch is trivially zero for any complete trajectory. v_2 is meaningful only on *prefixes* — partial paths of length k during generation — where it measures how many refinement steps the partial boundary condition survives. It is a *dynamics* measurement, not a topology measurement.

The corrected relationship

	Before	After	Why it changed
$\mathbb{Z}_5$	"True obstruction detector"	Homology with cyclic structure; global classifier	H1 is a class, not a detection signal
$\mathbb{Z}_2$	Weaker version of $\mathbb{Z}_5$	Orientation-forgetting projection; different invariant	Reversal experiment showed identity under $\mathbb{Z}_2$
2-adic	Part of topology	Prefix stability; generation-time dynamics	Full cycles have trivial boundary mismatch
Relationship	Hierarchy of detection power	Different invariants of one complex; not comparable	Experiments showed they answer different questions

The one-sentence correction

$\mathbb{Z}_2$ and $\mathbb{Z}_5$ are not different levels of hallucination detection — they are different projections of the same semantic chain complex: $\mathbb{Z}_5$ encodes global homology with orientation, $\mathbb{Z}_2$ is the orientation-forgetting quotient, and neither is a detector of truth — only of different invariants of a shared structure.

Why this matters for publication

The corrected framing is *stronger*, not weaker, for the following reason. The original framing ("different arithmetic detects different hallucinations") invites the question: why those particular number systems? The corrected framing answers it precisely: $\mathbb{Z}_2$ and $\mathbb{Z}_5$ are natural because they expose different algebraic invariants of the same chain complex — one orientation-sensitive, one orientation-free. 2-adic valuation is natural because generation is an iterative process and prefix stability is the right dynamical measurement. The three layers are not competing; they are complementary projections.

18.1 False positives: why over-restriction is the correct failure mode

Table 1 identifies this architecture's failure mode as *false positives* (overly restrictive admissibility) rather than false negatives (missed hallucinations). Reviewers may ask: does pruning novel but un-liftable trajectories harm generation diversity?

The answer depends on the deployment context, and for the systems this architecture targets, pruning is the correct failure mode.

High-SLA systems prefer false positives. In transit billing, medical routing, protocol state machines, and financial workflow automation:

- A false negative (permitting a structurally invalid continuation) can cause an unsafe execution loop, a billing anomaly, a mis-routed medical order, or a protocol violation with real downstream consequences.
- A false positive (pruning a novel but un-liftable trajectory) causes the system to decline and request clarification — recoverable, auditable, safe.

The asymmetry is fundamental: false negatives in these systems are *silent errors*; false positives are *visible refusals*. Regulators and auditors require the latter.

Diversity within the admissible support is preserved. The gate does not collapse generation to a single output. Within the admissible support $V(I)$, the full generative distribution P_θ remains active — probability determines selection, topology determines the boundary. Empirically: the gate reduces the branching factor to $0.50\times$ (50% candidate pruning) while preserving generation diversity within the admissible subspace.

The open-world extension. When the admissibility structure must accommodate genuinely novel valid combinations not present in the compiled topology, Pass 6 (cone compiler, cluster mutations) provides the mechanism: wall-crossing at the boundary of $V(I)$ is handled algebraically, not by relaxing the gate. The compiled structure grows; the gate remains hard.

18.2 Scope: where the architecture is strongest

The framework is strongest for domains that naturally admit continuation invariants, obstruction structure, and admissibility topology:

- **Protocol and workflow generation** (Phase 1 benchmark): state-machine transitions have exact admissibility topology.
- **Structured reasoning and theorem continuation:** proof steps have causal and logical continuation constraints.
- **Planning and routing** (MTR benchmark): network topology defines toric ideal generators exactly.
- **Program synthesis:** type systems and control-flow graphs provide natural Postnikov filtrations.
- **Stateful and causal trajectory generation:** any domain where “what states can follow what states” is definable in advance.

The architecture is *weaker* for open-ended natural language generation where admissibility topology cannot be pre-compiled — the false-positive rate (over-restriction) would be high. The correct deployment sequence is: Phase 1 (symbolic/structured) $\rightarrow$ Phase 2 (semi-structured reasoning) $\rightarrow$ Phase 3 (dynamic contextual topology) $\rightarrow$ Phase 4 (natural language) only after the above.

18.3 The central remaining empirical question

The key challenge identified in peer review is:

Do the admissibility constraints generalise, or do they merely memorise finite valid trajectories?

This is now the scientific heart of the experimental programme. The required experiments are:

- (i) **Compositional generalisation:** generate trajectories combining sub-sequences not seen in training. Does validity hold? Does the topology contribute nontrivially over a learned mask?
- (ii) **OOD continuation validity:** the temperature sweep (Phase 1 results: baseline 54% at $T = 3.0$, constrained 100%) is a start. Full OOD requires topology changes (new states, new edges) not seen during compilation.
- (iii) **Diversity preservation:** does the admissibility projection preserve the entropy of P_θ within the admissible support, or does it catastrophically prune? Measurement: generation diversity (self-BLEU, unique n -gram rate) inside and outside the gate.
- (iv) **Creativity preservation:** are valid but “creative” trajectories (low-probability but admissible) still reachable? This tests whether the gate is merely a memorised rejection table or a genuine structural filter.

18.4 Computational cost: compile-time vs runtime

The critical framing: all expensive algebra is compiled offline. At inference, the gate executes only sparse table lookups.

Compile-time (one-time, offline). The toric ideal generators, Markov basis, chain complex assembly, and admissibility table are computed once before deployment. On the 8-node protocol benchmark: total compile time = 0.97 ms. On the 81-station MTR network (37 Markov circuits, 161 Graver elements): compile time is dominated by 4ti2, a one-time computation amortized over all tokens served by the system.

Runtime gate cost (per-token).

Gate	Validity	ms/token	Overhead	Branching factor
No gate (baseline)	82.4%	0.058 ms	—	1.000×
FSM (P4 only)	100.0%	0.101 ms	+75%	0.500×
$\mathbb{Z}_2$ structural	95.0%	0.074 ms	+28%	0.625×
$\mathbb{Z}_5$ homology	100.0%	0.087 ms	+50%	0.500×
Full gate	100.0%	0.143 ms	+147%	0.500×

Runtime scales linearly. Gate cost scales linearly with vocabulary size (confirmed experimentally on graphs from 10 to 1000 nodes): doubling the graph doubles the gate cost, not the token generation cost. At inference the gate is a sparse loop over candidates — no Gröbner basis, no symbolic homology, no Floer solve.

Branching factor reduction. The gate reduces the effective branching factor to $0.50\times$ — 50% of candidate tokens are pruned before the softmax. This compresses downstream beam-search and reranking computation: fewer invalid candidates enter the probability simplex, reducing wasted gradient and sampling compute.

Correct framing for reviewers. Do not state: "we compute Floer cohomology online." State: "the admissibility structure is compiled offline into sparse runtime verification operators whose evaluation cost is comparable to standard FSM-constrained decoding."

18.5 Calibration: the tightened cohomological label rule

Running the Reverse Hironaka functor R_H against transformer hidden states revealed a systematic false-positive pattern: factual trajectories with high $v_{\text{fake}} > 1.5$ were being labelled HALLUCINATION despite exhibiting no topological failure. The root cause is the *anisotropic spectrum* (the "cone effect") of transformer latent representations: cross-window covariance spikes during semantic advection—when the model transitions between distinct contextual registers—cause transient v_{fake} elevations that reflect benign geometric distortion, not structural failure.

The old rule (over-restrictive).

$$\text{is_hallucination} \leftarrow v_{\text{fake}} > 1.5 \quad (1)$$

This scalar threshold treats v_{fake} as a topological invariant, which it is not. v_{fake} measures the p -adic absolute value of the *chain-level* Stasheff residual:

$$v_{\text{fake}}(x) = \|m_1^2(x) + m_2(m_0, x) + m_2(x, m_0)\|_p$$

This is a *continuous metric distance* to the zero-locus $V(I)$ of the toric ideal—not the cohomology class, which jumps discretely. High-variance language can produce $v_{\text{fake}} > 1.5$ while the underlying trajectory lifts cleanly through the $\mathbb{Z}_p$ tower.

The tightened rule (conjunction required).

$$\text{is_hallucination} \leftarrow (v_{\text{fake}} > 1.5) \wedge (\chi_{\text{gate}} < 0 \vee \text{Sq}^1(x) \neq 0) \quad (2)$$

where $\text{Sq}^1(x) \neq 0$ denotes the first Bockstein obstruction: failure to lift from $\mathbb{Z}/p\mathbb{Z}$ to $\mathbb{Z}/p^2\mathbb{Z}$.

The conjunction establishes an explicit decision diagram:

$$\begin{aligned} v_{\text{fake}} > 1.5, \chi_{\text{gate}} \geq 0, \text{Sq}^1 = 0 &\implies \text{BENIGN DISTORTION (Admitted)} \\ v_{\text{fake}} > 1.5, (\chi_{\text{gate}} < 0 \vee \text{Sq}^1 \neq 0) &\implies \text{HALLUCINATION (Blocked)} \end{aligned}$$

Why the conjunction is theoretically grounded. The rule rests on a strict separation of roles:

- v_{fake} : *metric distance* (continuous). Measures how far the trajectory is from $V(I)$ at the chain level. Can spike during benign semantic advection through anisotropic latent-space regions (the cone effect).
- $\chi_{\text{gate}} \in \{-1, 0, +1\}$: *topological invariant* (discrete). The Hodge–Euler decision $\text{sign}(f_{\text{grad}} - f_{\text{curl}}) \cdot \mathbf{1}_{\{\tilde{H}^1=0\}}$ from the sheaf complex. $\chi_{\text{gate}} = -1$ (DEAD) signals zero harmonic component: the trajectory lies on a dangling space with no global section.
- $\text{Sq}^1(x) \neq 0$: *arithmetic invariant* (discrete). Failure to lift through the first level of the $\mathbb{Z}_p$ tower $\cdots \rightarrow \mathbb{Z}/p^3 \rightarrow \mathbb{Z}/p^2 \rightarrow \mathbb{Z}/p \rightarrow 0$. When $\text{Sq}^1(x) \neq 0$, the inverse limit is undefined at x : the class lies on a dangling space with no $\mathbb{Z}_p$ -section.

High v_{fake} with $\chi_{\text{gate}} \geq 0$ and $\text{Sq}^1 = 0$ means the path is metrically distorted but topologically and arithmetically sound: admitted as UNCERTAIN, not blocked as HALLUCINATION. This is the semantic advection case.

Calibration results.

Trajectory type	HALLU	ADM	UNC	Interpretation
Factual ($1/k$ spectrum, layernorm)	0	0	7	All windows: benign distortion, no IL failure
Citation hallucination (drift + 7-adic cycle)	0	2	5	X^3, X^4 rescued (peak coherence window)
Sheaf hallucination (per-window rotation)	0	0	7	H^1 sheaf signal carries the detection

The factual trajectory returns **0 false positives**. The citation hallucination correctly identifies X^3 – X^4 (tokens 257–512) as ADMISSIBLE—the peak-coherence window before the 7-adic cycle begins—while surrounding windows receive logit warps in $[-8.0, -3.4]$.

Julia implementation (au_compiler.jl, Pass 7).

```
# Old rule (over-restrictive):
# is_hallucination = v_fake > 1.5

# Tightened cohomological rule (conjunction):
is_hallucination = (v_fake > 1.5) &&
(chi_gate < 0 || bockstein_failed_at_level_1)
```

18.6 Operator validation and F_{depth} failure

The F_{depth} metric defined as the first BSS page r at which the differencing operator $d_r = T^r - T^{r-1} \bmod p$ kills a class was tested against four known structured systems before being applied to hallucination detection. The result is unambiguous: $F_{\text{depth}} = 1$ **universally** across pure noise, induction head patterns, memorised content, and novel content (30 trials each, all primes $p \in \{2, 5, 7\}$, dimension 12).

Root cause. The p -adic lattice scaling maps the transition matrix T (with entries $\sim 10^{-1}$ to 10^{-2}) to $T_p = \lfloor T \cdot p/2 \|T\|_{\infty}^{-1} \rfloor \bmod p$, whose diagonal entries are mostly 0. Consequently $T_p - I \equiv -I \pmod{p}$, which has full rank n . Every cohomology class is killed at page 1 regardless of semantic content.

Spectral concentration metric. The replacement metric $\text{SC}(T) = \sum_i |\lambda_i|^2 / (\sum_i |\lambda_i|)^2$ (spectral concentration = inverse participation ratio) correctly identifies memorised content with $k = 5$ modes as $\text{SC} = 1/k = 0.25$, compared to $\text{SC} = 1/d \approx 0.083$ for isotropic noise. However, all hidden states from the synthetic structured model (PhasedGPT-2) give $\text{SC} = 1/d$ at all training phases, because the LS-fitted transition matrices have full rank d with uniform eigenvalue distribution — a consequence of the absence of semantic structure.

Conclusion. Neither F_{depth} nor spectral concentration discriminates factual from hallucinated text in an unstructured model. This is precisely the predicted failure mode: the metrics are dominated by the choice of operator T and discretization scheme rather than by the semantic content of the representations.

Prediction for trained models. Attention sinks in trained GPT-2 create stable low-rank subspaces in specific layers that would produce $SC \gg 1/d$ for factually grounded windows and $SC \approx 1/d$ for hallucinated windows (diffuse, no attention sink). The operator validation test must be re-run on trained model hidden states before these metrics can be applied to hallucination detection.

The only confirmed empirical finding. The H_1 persistence gap (HPI_{H_1}) from Rips complex persistent homology (computed via `ripser`) shows a sign flip at training phase 0.2–0.3 in the synthetic phase sweep: early phases (< 0.2): fabricated text has marginally more H_1 loops (noise-dominated); late phases (≥ 0.3): factual text develops organised H_1 structure exceeding the fabricated text (manifold contraction in hallucination). This is the only signal that does not go through F_{depth} , and it is currently the only genuine empirical result. Its confirmation requires trained model checkpoints with FActScore labels.

18.7 Known limitations

- (i) We operate at classical limit $T = 0$. Full quantum corrections (Novikov ring) constitute Passes 6–8.
- (ii) The guarantee applies at the topological boundary. Within $V(I)$, ranking is heuristic and requires empirical validation.
- (iii) The static routing table requires recompilation when ω changes substantially. Adapton-style incremental recompilation is identified but not implemented.
- (iv) The moduli problem is computed at leading order ($\#\mathcal{M} \approx \sqrt{L_i \cdot L_j \cdot \omega}$), not from full virtual fundamental chains.

19 Patent Claim Structure

The following structured claim captures the core invention in a form suitable for patent or priority disclosure:

Claim 1 (Method). A method for generative continuation validation comprising:

- (a) a compiled Postnikov filtration $\{P_1, \dots, P_5\}$ over an admissible semantic manifold (M, ω) , wherein each stage corresponds to a distinct semantic coherence scale (token, phrase, temporal, causal, discourse);
- (b) obstruction cocycles ob_n evaluated prior to token sampling via pre-compiled, memory-mapped constraint structures (toric ideal generators, Floer-like continuation invariants, HMM brackets);
- (c) a support restriction rule $P_{\text{topo}}(x) \propto P_\theta(x) \cdot \prod_{n=1}^5 \mathbf{1}_{\{ob_n=0\}}$ applied before the softmax normalisation; and
- (d) wherein a continuation is a *structural hallucination* iff $\exists n : ob_n \neq 0$, and such continuations receive exactly zero probability mass.

Claim 2 (System). A system comprising: a static offline compiler that ingests a network topology and emits pre-compiled admissibility tables; a runtime gate that consults these tables in $O(1)$ per token; and an autoregressive decoder whose sampling distribution has support restricted to trajectories passing all five Postnikov stages, requiring no external factual oracle at inference time.

20 What Is a Hallucination in This Framework?

After the algebraic corrections, the framework gives a precise, three-level answer that differs fundamentally from all prior definitions.

Not any of these

A hallucination is *not*:

- a low-probability token (standard uncertainty framing)
- a factually incorrect statement (requires external oracle)
- a high-entropy output (misses confident hallucinations)
- a token outside the training distribution (statistical framing)

But this

A hallucination is a trajectory that *fails admissibility in the semantic chain complex* — it violates structural consistency in at least one of three algebraic projections:

Layer	What it checks	Failure = hallucination type
$H_1(C_*; \mathbb{Z}_5)$	"Does the full reasoning path close correctly? Is the cycle homologous to a filled disk?"	Logical contradiction; globally inconsistent story; impossible world state; "sounds right locally but cannot complete"
$\mathbb{Z}_2$ projection	"Are we conflating two structurally different paths because orientation/sign is lost?"	Ambiguity collapse; symmetric-but-distinct paths merged; reversed reasoning treated as equivalent
2-adic v_2 (prefix)	"Does the trajectory start becoming inconsistent before it fully forms?"	Semantic drift; compounding error; early divergence from admissible structure

Operational definition

$$\begin{aligned} \text{hallucination} &\iff \neg(\text{stable prefix}) \\ &\vee \neg(\text{structurally distinct path}) \\ &\vee \neg(\text{globally fillable cycle}) \end{aligned}$$

Why this is stronger than prior definitions

1. **No external oracle.** The admissibility check uses only the compiled chain complex — no retrieval database, no human labels, no verifier model.

2. **Pre-sampling.** The gate fires before the softmax. Hallucinations are not detected after generation and corrected; they are structurally unreachable.
3. **Three failure modes, not one.** Prior work detects one type of failure per method. This framework decomposes failure into global (H_1), structural ($\mathbb{Z}_2$), and dynamical (2-adic) components.
4. **Invariant under smooth perturbation.** The homological admissibility check does not change when embeddings are perturbed, weights are rescaled, or sampling temperature is adjusted. It is a topological invariant, not a statistical one.

One-sentence version

In this framework, hallucinations are trajectories that fail admissibility in the semantic chain complex, where $\mathbb{Z}_5$ detects global non-fillability, $\mathbb{Z}_2$ detects structural collapse under orientation loss, and 2-adic valuation detects instability during prefix formation.

21 Conclusion

We have presented a framework in which hallucination removal is algebraic and topological rather than probabilistic. The key objects — toric ideal, Floer chain complex, Postnikov filtration, Hodge decomposition, AU sheaf — are each operationally defined and computationally instantiated on the 81-station MTR network.

The central claim is:

Hallucination is reformulated as a failure of hierarchical continuation liftability through a compiled semantic filtration. Obstruction cocycles are evaluated prior to sampling, restricting the support of the generative distribution to trajectories admitting coherent continuation across all active Postnikov stages. No external ground-truth oracle is required at inference time. Topology determines support; probability determines selection within support.

We are careful to maintain three precise distinctions:

- Geometric obstruction $\neq$ semantic falsity. A blocked trajectory is unsupported by the compiled topology, not necessarily false in all possible worlds.
- Binary gate $\neq$ deterministic system. Within the admissible support, P_θ remains probabilistic.
- Structural constraint $\neq$ omniscience. The admissibility system encodes the compiled topology, not universal human knowledge.

The claim is:

A compiled semantic admissibility architecture in which generation is transformed from pointwise likelihood maximisation into sampling only from liftable continuations — where liftability is tested prior to softmax normalisation, at multiple semantic scales, through compiled runtime admissibility projection structures, requiring no external factual oracle at inference time.

This is a nontrivial conceptual shift: standard transformers *maximise local likelihood*; this architecture *samples only from liftable continuations*. The empirical correlate is demonstrated in the Phase 1 benchmark (Table 2: 100% validity at all temperatures vs. baseline collapse to 54% at $T = 3.0$). The deepest remaining question — whether the topology generalises compositionally beyond the compiled training manifold — is the central programme of Phases 2–4.

One-line distilled version. The experiments suggest that hallucination-like failures decompose into: *local statistical plausibility* ($\mathbb{R}$, gradient component), *structural admissibility* ($\mathbb{Z}_2$, orientation-free), and *global compositional closure* ($\mathbb{Z}_5$, H_1) — each enforced by progressively stronger algebraic constraints with near-linear runtime scaling. The admissibility machinery scales approximately linearly while reducing branching factor by 50 %, compressing the generative search space rather than merely blocking tokens. The strongest systems result is not 100 % validity but this: *topological constraints enforce global closure with $O(|E|)$ inference overhead, the same asymptotic complexity as standard FSM decoding*.

The strongest formal statement is the Homological Admissibility Theorem 9.2: any generated sequence whose continuation induces a trivial obstruction class at any Postnikov stage is excluded from decoding support — algebraically, not statistically.

Future work: compiler passes

The immediate algebraic programme comprises:

- **Pass 6** (Cone compiler): wall-crossing cluster mutations at interchange stations via SyzygyProjection.
- **Pass 7** (Vanishing cycles): seasonal Dehn twist monodromy via ExceptionalProjection.
- **Pass 8** (Stalk transport): Gauss-Manin connection for full m_1 parallel transport via SpectralProjection.

What the current architecture catches and what it misses

The experiments established a precise scope.

Caught by the current gate.

- (i) **Illegal transitions** (P_1 , P_4): nonexistent station, impossible protocol state, invalid edge traversal.
- (ii) **Non-fillable loops** ($H_1(\mathbb{Z}_5)$): trajectory cannot close coherently; discourse path lacks continuation class; cycle bounds a hole, not a filled disk.
- (iii) **Orientation/refinement failures** (2-adic v_2): contradiction emerging only after long continuation depth; unstable prefix under binary refinement.

Not caught by the current gate.

Hallucination type	Example	Why current system misses it
Factual	"Paris is in Germany"	Topologically admissible; forms a valid chain
Semantic substitution	Wrong entity	Still closes coherently in H_1
Numerical	Wrong arithmetic	Closure unaffected by arithmetic errors
World-model	Fake citation	Trajectory remains graph-valid
Discourse drift	Topic subtly changes	No topological obstruction
Ontological mismatch	"A cat is a prime number"	May remain graph-valid

The fundamental reason: topology checks *"Can this continuation exist coherently?"* It does not check *"Is this continuation true in the external world?"* This distinction is now one of the strongest conceptual clarifications in the work.

Where the missing hallucinations live: sheaf gluing failures. The hallucinations that survive all current invariants are *globally inconsistent semantic assignments* — local sections that are each valid but incompatible at their overlaps. This is a sheaf-theoretic failure, not a topological one.

Window 1: "Alice lives in Boston."

Window 2: "Alice moved to Seattle."

Window 3: "Alice currently commutes daily from Boston."

Each window is topologically valid ($H_1 = 0$, no obstruction). But the global section is inconsistent: Alice cannot simultaneously live in Seattle and commute from Boston. This is not a $\mathbb{Z}_5$ obstruction. It is a *gluing failure* — the local sections do not agree on their overlaps. Formally: $\check{H}^1(\mathcal{U}, \mathcal{F}) \neq 0$ for the sheaf $\mathcal{F}$ of entity-state assignments.

The hierarchy.

Layer	Tool	Catches	Misses
1	$\mathbb{R}$ (softmax)	Local plausibility	Structural impossibility
2	$\mathbb{Z}_2$	Structural presence	Global closure
3	$\mathbb{Z}_5$ H_1	Non-fillable trajectories	Semantic truth
4	Sheaf $\check{H}^1$	Cross-context semantic gluing	External world facts

The stronger claim that emerges:

Hallucinations arise from either non-liftable continuations (topological obstruction, caught by $H_1(\mathbb{Z}_5)$) or incompatible semantic gluings (sheaf obstruction, requiring $\check{H}^1(\mathcal{U}, \mathcal{F})$). This architecture handles the first class completely; the second class is the natural next layer.

Future work: empirical validation requirements

The framework is beyond idea stage. To achieve publication, the following empirical evidence is required:

- (i) **Benchmark comparison:** baseline transformer vs. homologically-gated transformer on standard hallucination benchmarks (TruthfulQA, HaluEval, FactScore).
- (ii) **Invalid trajectory rate:** measure the fraction of generated sequences that are topologically blocked at each Postnikov stage. This quantifies how much work the gate is actually doing.
- (iii) **Long-context consistency:** demonstrate that the Floer-like continuation gate maintains coherence over longer contexts than the baseline.
- (iv) **OOD robustness:** show that the topology contributes nontrivially on out-of-distribution inputs — not merely on the training distribution.
- (v) **Ablation studies:** remove each Postnikov stage independently and measure the hallucination rate at that level. This validates the hierarchical structure.

- (vi) **Language model instantiation:** define the Lagrangians, symplectic form, and toric ideal for a concrete NLP task and validate the gate on factual QA.

Without these, reviewers will correctly characterise the system as “a symbolic validator attached to a transformer.” The experiments must prove that the topology contributes nontrivially — that the admissibility structure produces compositional generalisation and emergent long-range consistency that a standard learned mask cannot replicate.

References

- [1] K. Fukaya, Y.-G. Oh, H. Ohta, K. Ono, *Lagrangian Floer Theory and Mirror Symmetry on Compact Toric Manifolds*, arXiv:1009.1648v3 [math.SG], 292 pp., 2016.
- [2] B. Fang, C.-C. M. Liu, D. Treumann, E. Zaslow, *T-duality and Homological Mirror Symmetry for Toric Varieties*, Adv. Math. **229** (2012), 1873–1911.
- [3] P. Seidel, *Fukaya Categories and Picard-Lefschetz Theory*, European Mathematical Society, 2008.
- [4] 4ti2 team, *4ti2 — A software package for algebraic, geometric and combinatorial problems on linear spaces*, <https://4ti2.github.io>.
- [5] P. Seidel, *A long exact sequence for symplectic Floer cohomology*, arXiv:math/0105186v2 [math.SG], revised May 2002.
Exact results used: **Theorem 1** — Long exact sequence $HF(\tau_L(L_0), L_1) \rightarrow HF(L_0, L_1) \rightarrow HF(L, L_1) \otimes HF(L_0, L)$; used in §12 to ground the Dehn twist framing. **Proposition 1.15** — Monodromy isotopic to Dehn twist along vanishing cycle; used to explain hallucination stability under prompt perturbation.
- [6] P. Seidel, *Localization for involutions in Floer cohomology*, *Geom. Funct. Anal.* **20**(6):1464–1501, 2010. arXiv:0907.2748 [math.SG].
Theorem 1: Smith inequality for Floer cohomology. The localization map $\Lambda(m)$ becomes an isomorphism after inverting the Novikov variable q ; the $\mathbb{Z}_5$ valuation gate in this paper plays the same role: after passing to $v_5 \geq \tau$, the obstruction that was invisible in $\mathbb{Z}_2$ becomes detectable.
- [7] P. Seidel, *Some speculations on pairs-of-pants decompositions and Fukaya categories*, *Proceedings of Gökova Geometry-Topology Conference 2010*, pp. 1–13, 2011. arXiv:1004.0906 [math.SG].
Conjecture 2.2: $\mathcal{F}(M) \hookrightarrow DC^*(W)$. Pairs-of-pants = P_1 – P_4 atomic local checks; global gluing = P_5 sheaf condition. The Novikov field Λ_q with valuation val_q (Section 3, Eq. (3.1)) is the theoretical basis for the 5-adic valuation gate; the Legendre transform ψ and toric degeneration $W : \mathcal{X} \rightarrow \mathbb{A}^1$ correspond to the compiled toric ideal and routing table.
- [8] P. Seidel, *Fukaya categories and deformations*, *Proc. ICM (Beijing)* **2**:351–360, 2002. arXiv:math/0206155 [math.SG].
Open-closed string map $SH^*(M) \rightarrow HH^*(\mathcal{F}(M), \mathcal{F}(M))$; A_∞ deformation parameter t = Novikov variable = 5-adic valuation depth in this paper. An unobstructed deformation ($\mu_0 = 0$) corresponds to $ob_n = 0$ throughout the Postnikov tower.
- [9] Y.-G. Oh, *Floer cohomology of Lagrangian intersections and pseudo-holomorphic disks I*, *Commun. Pure Appl. Math.* **46**(7):949–994, 1993.
- [10] C. Viterbo, *Functors and computations in Floer homology with applications, Part I*, *Geom. Funct. Anal.* **9**(5):985–1033, 1999.